



**RAMNIRANJAN JHUNJHUNWALA COLLEGE
GHATKOPAR (W), MUMBAI - 400 086**

**DEPARTMENT OF
INFORMATION TECHNOLOGY**

2021 - 2022

M.Sc.(I.T.) Part II (Sem - 3)

Subject : Deep Learning

**Name : Sneha Ramchandra Pawar. Roll
No. : 17**



Hindi Vidya Prachar Samiti's

RAMNIRANJAN

**JHUNJHUNWALA COLLEGE
(AUTONOMOUS)**



Opposite Ghatkopar Railway Station, Ghatkopar West, Mumbai-400086

CERTIFICATE

This is to certify that Miss. **Sneha Ramchandra Pawar** with Roll No. **17** has successfully completed the necessary course of experiments in the subject of Deep **Learning** during the academic year **2021 – 2022** complying with the requirements of **RAMNIRANJAN JHUNJHUNWALA COLLEGE OF ARTS, SCIENCE AND COMMERCE**, for the course of **M.Sc. (IT)** semester -I.

Internal Examiner

External Examiner

Head Of Department

College Seal

Index

Sr. No.	Title	Date	Page No.
1	Matrix Multiplication, Eigen Vectors, EigenValue Computation using TensorFlow		
2	Deep Forward Network for XOR		
3a	Classification using DNN		
3b	Binary Classification using MLP		
3c	Convolutional Neural Network		
4a	Feed Forward NN		
4b	Predicting the Probability of the class		
5a	CNN for CIFAR10 Images		
5b	Image Classification		
5c	Data Augmentation		
6	Building RNN using Single Neuron		
7	Using CovNet to build Deep Learning model		
8	Implementing a single Autoencoder based on fully connected layer		

Practical No :- 1

Aim:-Matrix Multiplication, Eigen Vectors, EigenValue Computation using TensorFlow

Matrix Multiplication:-

In mathematics, particularly in linear algebra, **matrix multiplication** is a binary operation that produces a matrix from two matrices. For matrix multiplication, the number of columns in the first matrix must be equal to the number of rows in the second matrix.

Eigen Vectors:-

Eigenvectors are the vectors (non-zero) that do not change the direction when any linear transformation is applied. It changes by only a scalar factor. In a brief, we can say, if A is a linear transformation from a vector space V and x is a vector in V, which is not a zero vector, then v is an eigenvector of A if $A(x)$ is a scalar multiple of x .

An **Eigenspace** of vector x consists of a set of all eigenvectors with the equivalent eigenvalue collectively with the zero vector. Though, the zero vector is not an eigenvector.

Let us say A is an “ $n \times n$ ” matrix and λ is an eigenvalue of matrix A, then x , a non-zero vector, is called as eigenvector if it satisfies the given below expression;

$$Ax = \lambda x$$

x is an eigenvector of A corresponding to eigenvalue, λ .

EigenValue:-

Eigenvalues are the special set of scalars associated with the system of linear equations. It is mostly used in matrix equations. ‘Eigen’ is a German word that means ‘proper’ or ‘characteristic’. Therefore, the term eigenvalue can be termed as characteristic value, characteristic root, proper values or latent roots as well. In simple words, the eigenvalue is a scalar that is used to transform the eigenvector. The basic equation is

$$Ax = \lambda x$$

The number or scalar value “ λ ” is an eigenvalue of A.

In Mathematics, an eigenvector corresponds to the real non zero eigenvalues which point in the direction stretched by the transformation whereas eigenvalue is considered as a factor by which it is stretched. In case, if the eigenvalue is negative, the direction of the transformation is negative.

For every real matrix, there is an eigenvalue. Sometimes it might be complex. The existence of the eigenvalue for the complex matrices is equal to the fundamental theorem of algebra.

What is TensorFlow?

TensorFlow is an open-source end-to-end platform for creating Machine Learning applications. It is a symbolic math library that uses dataflow and differentiable programming to perform various tasks focused on training and inference of deep neural networks. It allows developers to create machine learning applications using various tools, libraries, and community resources. Currently, the most famous deep learning library in the world is Google's TensorFlow. Google product uses machine learning in all of its products to improve the search engine, translation, image captioning or recommendations.

<https://colab.research.google.com/drive/12PJdyaP5mhK0E11NMTKogyNW2YHrp8wK>

Code:-

```
import tensorflow as tf
print("Matrix Multiplication Demo")
x=tf.constant([1,2,3,4,5,6],shape=[2,3])
print(x)
y=tf.constant([7,8,9,10,11,12],shape=[3,2])
print(y)
z=tf.matmul(x,y)
print("Product:",z)

e_matrix_A=tf.random.uniform([2,2],minval=3,maxval=10,dtype=tf.float32,name="matrixA")

print("Matrix A:\n{}\n\n".format(e_matrix_A))

eigen_values_A,eigen_vectors_A=tf.linalg.eigh(e_matrix_A)

print("Eigen Vectors:\n{}\n\nEigen values:\n{}\n".format(eigen_vectors_A, eigen_values_A))
```

Output:-

```

import tensorflow as tf
print("Matrix Multiplication Demo")
x=tf.constant([1,2,3,4,5,6],shape=[2,3])
print(x)
y=tf.constant([7,8,9,10,11,12],shape=[3,2])
print(y)
z=tf.matmul(x,y)
print("Product:",z)

e_matrix_A=tf.random.uniform([2,2],minval=3,maxval=10,dtype=tf.float32,name="matrixA")
e="matrixA"

print("Matrix A:\n{}\n{}\n".format(e_matrix_A))

eigen_values_A,eigen_vectors_A=tf.linalg.eigh(e_matrix_A)

print("Eigen Vectors:\n{}\n{}\nEigen values:\n{}\n{}\n".format(eigen_vectors_A, eigen_values_A))

```

Matrix Multiplication Demo

```

tf.Tensor(
[[1 2 3]
 [4 5 6]], shape=(2, 3), dtype=int32)
+ tf.Tensor(
[[ 7  8]
 [ 9 10]
 [11 12]], shape=(3, 2), dtype=int32)
Product: tf.Tensor(
[[ 58  64]
 [139 154]], shape=(2, 2), dtype=int32)
Matrix A:
[[8.621217  5.0865545]
 [3.6364348 6.520332 ]]

Eigen Vectors:
[[[-0.60103273  0.79922444]
 [ 0.79922444  0.60103273]]]

Eigen values:
[ 3.7856605 11.355889 ]

```

Practical No :- 2

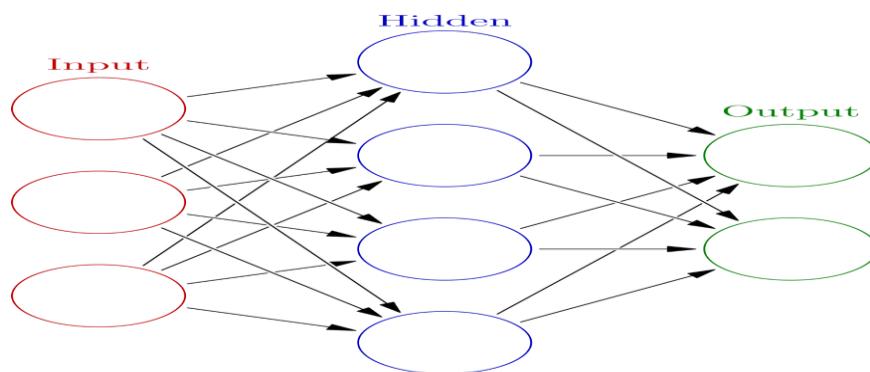
Aim:- Deep Forward Network for XOR

Deep feedforward networks, also often called feedforward neural networks, or multilayer perceptrons (MLPs), are the quintessential deep learning models.

These models are called feedforward because information flows through the function being evaluated from x , through the intermediate computations used to define f , and finally to the output y . There are no feedback connections in which outputs of the model are fed back into itself. When feedforward neural networks are extended to include feedback connections, they are called recurrent neural networks,

Feedforward networks are of extreme importance to machine learning practitioners. They form the basis of many important commercial applications. For example, the convolutional networks used for object recognition from photos are a specialized kind of feedforward network. Feedforward networks are a conceptual stepping stone on the path to recurrent networks, which power many natural language applications.

Feedforward neural networks are called networks because they are typically represented by composing together many different functions. The model is associated with a directed acyclic graph describing how the functions are composed together. For example, we might have three functions $f(1)$, $f(2)$, and $f(3)$ connected in a chain, to form $f(x) = f(3)(f(2)(f(1)(x)))$. These chain structures are the most commonly used structures of neural networks. In this case, $f(1)$ is called the first layer of the network, $f(2)$ is called the second layer, and so on.



<https://colab.research.google.com/drive/1d7tosUst7kqIKzP94YaJdH2qhwWXnEm4>

Code:-

```

import numpy as np
from keras.layers import Dense
from keras.models import Sequential
model=Sequential()
model.add(Dense(units=2,activation='relu',input_dim=2))
model.add(Dense(units=1,activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
print(model.get_weights())
X=np.array([[0.,0.],[0.,1.],[1.,0.],[1.,1.]])
Y=np.array([0.,1.,1.,0.])
model.fit(X,Y,epochs=1000,batch_size=4)
print(model.get_weights())
print(model.predict(X,batch_size=4))

```

Output:-

The screenshot shows a Google Colab notebook titled "Practical No:2.ipynb". The code cell contains the following Python code:

```

[1] #Practical No:2
#Aim: Solving XOR problem using deep feed forward network.

import numpy as np
from keras.layers import Dense
from keras.models import Sequential

model=Sequential()
model.add(Dense(units=2,activation='relu',input_dim=2))
model.add(Dense(units=1,activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())

```

The code cell has been executed successfully, indicated by the green checkmark icon. The output shows the model summary:

```

Model: <keras.Sequential>
_________________________________________________________________
Layer (type)                 Output Shape              Param #   
=================================================================
dense_1 (Dense)              (None, 2)                6        
dense_2 (Dense)              (None, 1)                3        
Total params: 9
Trainable params: 9
Non-trainable params: 0
_________________________________________________________________

```

The status bar at the bottom right of the Colab interface shows "0s completed at 21:28" and the date "26-11-2021".

My Drive - Google Drive Practical No2.ipynb - Colaboratory Practical 1 to 3

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[4] print(model.get_weights())
```

<> [array([[1.2027003 , -0.9790218], [-1.0076371 , 0.92028206]], dtype=float32), array([0., 0.], dtype=float32), array([-0.87117094, -0.30598032]], dtype=float32), array([0., 0.], dtype=float32)] {x}

Q X=np.array([[0.,0.],[0.,1.],[1.,0.],[1.,1.]]) Y=np.array([0.,1.,1.,0.])

model.fit(X,Y,epochs=1000,batch_size=4)

```
print(model.get_weights())
```

```
print(model.predict(X,batch_size=4))
```

Epoch 130/1000
1/1 [=====] - 0s 7ms/step - loss: 0.8043 - accuracy: 0.0000e+00
Epoch 131/1000
1/1 [=====] - 0s 7ms/step - loss: 0.8039 - accuracy: 0.0000e+00

Executing (2s) C_> error_hand... > fi... > on_train_batch_... > _call_batch_h... > _call_batch_end_h... > _call_batch_hook_h... > on_train_batch_... > _batch_update_pro... > upda... > wri... > schedu... 2130 29°C Haze 26-11-2021

Type here to search

My Drive - Google Drive Practical No2.ipynb - Colaboratory Practical 1 to 3

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
Epoch 992/1000  
1/1 [=====] - 0s 4ms/step - loss: 0.6118 - accuracy: 0.7500  
Epoch 993/1000  
1/1 [=====] - 0s 17ms/step - loss: 0.6116 - accuracy: 0.7500  
Epoch 994/1000  
1/1 [=====] - 0s 12ms/step - loss: 0.6114 - accuracy: 0.7500  
Epoch 995/1000  
1/1 [=====] - 0s 17ms/step - loss: 0.6111 - accuracy: 0.7500  
Epoch 996/1000  
1/1 [=====] - 0s 10ms/step - loss: 0.6109 - accuracy: 0.7500  
Epoch 997/1000  
1/1 [=====] - 0s 10ms/step - loss: 0.6107 - accuracy: 0.7500  
Epoch 998/1000  
1/1 [=====] - 0s 10ms/step - loss: 0.6105 - accuracy: 0.7500  
Epoch 999/1000  
1/1 [=====] - 0s 5ms/step - loss: 0.6102 - accuracy: 0.7500  
Epoch 1000/1000  
1/1 [=====] - 0s 5ms/step - loss: 0.6100 - accuracy: 0.7500
```

[array([[0.83197623, -1.3286294], [-0.2738633 , 1.3283621]], dtype=float32), array([-0.3707247 , -0.00040202], dtype=float32), array([-0.19533268, 0.66419214]), dtype=float32), array([-0.12735762], dtype=float32)]
[[0.46820357], [0.6801939], [0.44584945], [0.4591014]]

21s completed at 21:31 21:31 29°C Haze 26-11-2021

Type here to search

Practical No :- 3

Aim:- 3a.Classification using DNN

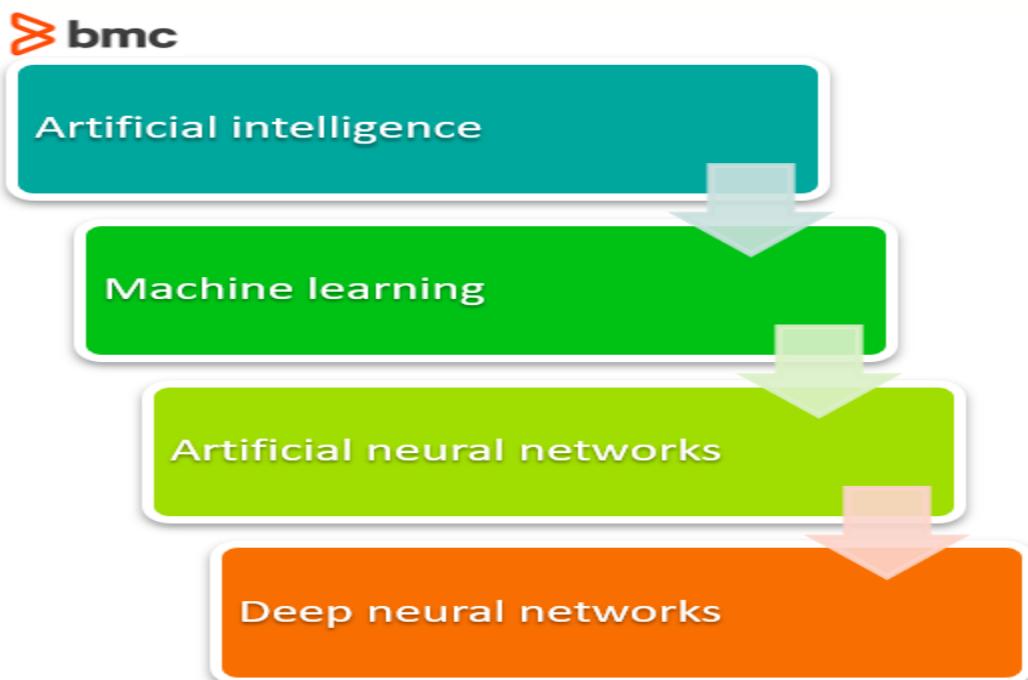
<https://colab.research.google.com/drive/19seC-t6aJplfEDSb2F7JVy2acVAUJeaH>

<https://colab.research.google.com/drive/1QJcZT9FU2RQvBGGzx1O6SUWQX1I2MJxn>

What is a deep neural network?

At its simplest, a neural network with some level of complexity, usually at least two layers, qualifies as a deep neural network (DNN), or deep net for short. Deep nets process data in complex ways by employing sophisticated math modeling.

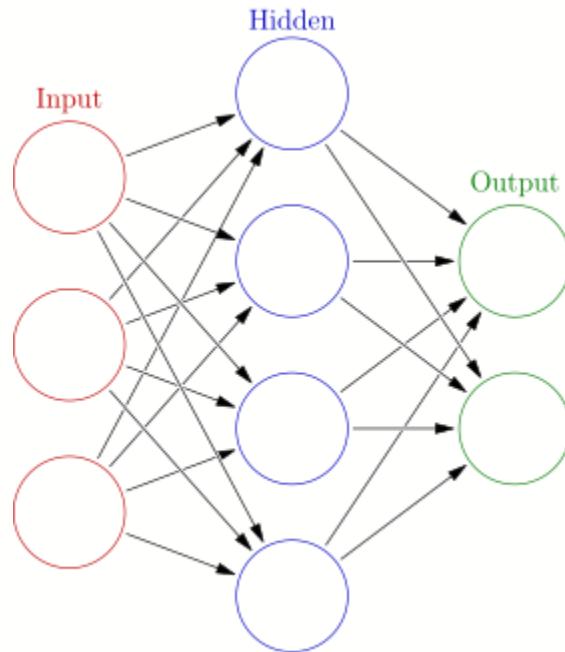
To truly understand deep neural networks, however, it's best to see it as an evolution. A few items had to be built before deep nets existed.



First, machine learning had to get developed. ML is a framework to automate (through algorithms) statistical models, like a linear regression model, to get better at making predictions. A model is a single model that makes predictions about something. Those predictions are made with some accuracy. A model that learns—machine learning—takes all its bad predictions and tweaks the weights inside the model to create a model that makes fewer mistakes.

The learning portion of creating models spawned the development of artificial neural networks. ANNs utilize the hidden layer as a place to store and evaluate how significant one of the inputs

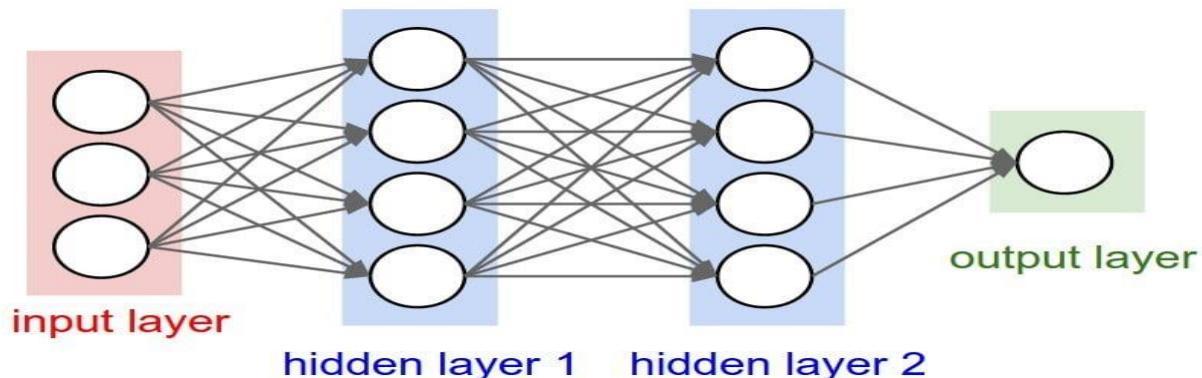
is to the output. The hidden layer stores information regarding the input's importance, and it also makes associations between the importance of combinations of inputs.



One hidden layer is considered an Artificial Neural Network (ANN)

Deep neural nets, then, capitalize on the ANN component. They say, if that works so well at improving a model—because each node in the hidden layer makes both associations and grades importance of the input to determining the output—then why not stack more and more of these upon each other and benefit even more from the hidden layer?

So, the deep net has multiple hidden layers. ‘Deep’ refers to a model’s layers being multiple layers deep.



Two or more hidden layers comprise a Deep Neural Network

Code-:

```
from numpy import loadtxt
```

```

from keras.models import Sequential
from keras.layers import Dense
import numpy
dataset = numpy.loadtxt("/content/sample_data/pima-indians-diabetes - pima-indians-diabetes.csv",
delimiter=",")
dataset
X = dataset[:,0:8]
Y = dataset[:,8]
X

Y
model = Sequential()
model.add(Dense(12, input_dim=8, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

model.fit(X, Y, epochs=150, batch_size=10)

#predictions = model.predict(X)

#rounded = [round(x[0]) for x in predictions]
#print(rounded)

#scores = model.evaluate(X, Y)
#print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
_,accuracy=model.evaluate(X,Y)
print('accuracy of model is',(accuracy*100) )
predictions = model.predict(X)

```

Output-:

```

[ ] #Aim: Implementing a deep neural network for performing classification task.
#Problem statement:
#The given dataset comprises health information about diabetic women patients. We need to create a deep
#neural network which can predict if a person has diabetes or not based on the input parameters.

from numpy import loadtxt
from keras.models import Sequential
from keras.layers import Dense
import numpy

dataset = numpy.loadtxt("/content/pima-indians-diabetes.csv", delimiter=",")

X = dataset[:,0:8]
Y = dataset[:,8]

model = Sequential()
model.add(Dense(12, input_dim=8, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

model.fit(X, Y, epochs=150, batch_size=10)

```

Executing (8s) Cell > error_handler() > fit() > steps() > numpy() > read_value() > error_handler() > op_dispatch_handler() > identity() > identity()

```

My Drive - Google Driv... | Common Doc - Goog... | Common Doc | pima-indians-diab... | Practical No2.ipynb | Practical 1 to 3 | practical3.ipynb - Col...
File Edit View Insert Runtime Tools Help Cannot save changes
+ Code + Text Copy to Drive
RAM Disk Editing
0s
21s
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

<>
model.fit(X, Y, epochs=150, batch_size=10)

{x}

#predictions = model.predict(X)

#rounded = [round(x[0]) for x in predictions]
#print(rounded)

#scores = model.evaluate(X, Y)
#print("\n%: %.2f%%" % (model.metrics_names[1], scores[1]*100))

Epoch 1/150
77/77 [=====] - 1s 2ms/step - loss: 13.8000 - accuracy: 0.4479
Epoch 2/150
0s completed at 21:40

```

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

0s completed at 21:40

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

0s completed at 21:43

Aim:-3b.Binary Classification using MLP

Code:-

```

# mlp for binary classification
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
# load the dataset
path ='/content/sample_data/Ionosphere - Ionosphere.csv'
df = read_csv(path, header=None)
# split into input and output columns
X, y = df.values[:, :-1], df.values[:, -1]
# ensure all data are floating point values
X = X.astype('float32')
# encode strings to integer
y = LabelEncoder().fit_transform(y)
# split into train and test datasets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
# determine the number of input features
n_features = X_train.shape[1]
# define model
model = Sequential()
model.add(Dense(10, activation='relu', kernel_initializer='he_normal', input_shape=(n_features,)))
model.add(Dense(8, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(1, activation='sigmoid'))
# compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
# fit the model
history=model.fit(X_train, y_train, epochs=150, batch_size=32, verbose=0)
# evaluate the model
loss, acc = model.evaluate(X_test, y_test, verbose=0)
print('Test Accuracy: %.3f' % acc)
# make a prediction
row = [1.0, 0.99539, -0.05889, 0.85243, 0.02306, 0.83398, -0.37708, 1.0, 0.03760, 0.85243, -0.17755, 0.59755, -0.44945, 0.60536, -0.38223, 0.84356, -0.38542, 0.58212, -0.32192, 0.56971, -0.29674, 0.36946, 0.47357, 0.56811, -0.51171, 0.41078, -0.46168, 0.21266, -0.34090, 0.42267, -0.54487, 0.18641, -0.45300]
yhat = model.predict([row])
print('Predicted: %.3f' % yhat)

from matplotlib import pyplot
#plot learning curves
pyplot.title('Learning Curves')
pyplot.xlabel('Epoch')
pyplot.ylabel('Cross Entropy')
pyplot.plot(history.history['loss'], label='train')
#pyplot.plot(history.history['val_loss'],label='val')
pyplot.legend()
pyplot.show()

```

Output:-

```
# determine the number of input features
n_features = X_train.shape[1]
# define model
model = Sequential()
model.add(Dense(10, activation='relu', kernel_initializer='he_normal', input_shape=(n_features,)))
model.add(Dense(8, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(1, activation='sigmoid'))
# compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
# fit the model
history=model.fit(X_train, y_train, epochs=150, batch_size=32, verbose=0)
# evaluate the model
loss, acc = model.evaluate(X_test, y_test, verbose=0)
print('Test Accuracy: %.3f' % acc)
# make a prediction
row = [1,0,0.99539,-0.05889,0.85243,0.02306,0.83398,-0.37708,1,0.03760,0.85243,-0.17755,0.59755,-0.44945,0.60536,-0.38223,0.84356,-0.38542,0.58
yhat = model.predict([row])
print('Predicted: %.3f' % yhat)

(235, 34) (116, 34) (235,) (116,
Test Accuracy: 0.914
Predicted: 0.993
```



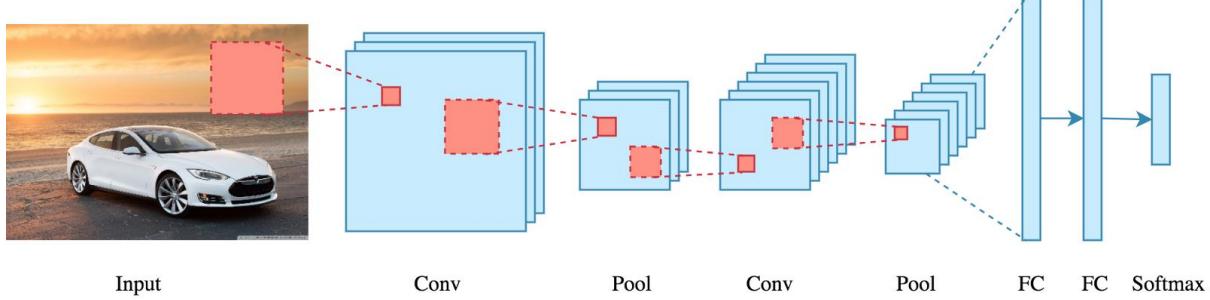
```
# DETERMINE MODEL
model = Sequential()
model.add(Dense(10, activation='relu', kernel_initializer='he_normal', input_shape=(n_features,)))
model.add(Dense(8, activation='relu', kernel_initializer='he_normal'))
model.add(Dense(1, activation='sigmoid'))
# compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
# fit the model
history=model.fit(X_train, y_train, epochs=150, batch_size=32, verbose=0)
# evaluate the model
loss, acc = model.evaluate(X_test, y_test, verbose=0)
print('Test Accuracy: %.3f' % acc)
# make a prediction
row = [1,0,0.99539,-0.05889,0.85243,0.02306,0.83398,-0.37708,1,0.03760,0.85243,-0.17755,0.59755,-0.44945,0.60536,-0.38223,0.84356,-0.38542,0.58
yhat = model.predict([row])
print('Predicted: %.3f' % yhat)

(235, 34) (116, 34) (235,) (116,
Test Accuracy: 0.940
Predicted: 0.998
```

Aim-:3c.Convolutional Neural Network

What exactly is a CNN?

In deep learning, a convolutional neural network (CNN/ConvNet) is a class of deep neural networks, most commonly applied to analyze visual imagery. Now when we think of a neural network we think about matrix multiplications but that is not the case with ConvNet. It uses a special technique called Convolution. Now in mathematics convolution is a mathematical operation on two functions that produces a third function that expresses how the shape of one is modified by the other.



But we don't really need to go behind the mathematics part to understand what a CNN is or how it works.

Bottom line is that the role of the ConvNet is to reduce the images into a form that is easier to process, without losing features that are critical for getting a good prediction.

Code:-

```
#Import TensorFlow
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0
#Verify the data
#To verify that the dataset looks correct, let's plot the first 25 images from the training set and display the class name below each image:
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)

    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    # The CIFAR labels happen to be arrays,
    # which is why you need the extra index
    plt.xlabel(class_names[train_labels[i][0]])

plt.show()
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

model.summary()
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))
model.summary()
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
```

```
        validation_data=(test_images, test_labels))
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(test_acc)
```

Output-:

practical 3(b,c).ipynb

```
#practical 3(c) Convolutional Neural Network (CNN)
import tensorflow as tf

from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 2s 0us/step
170508288/170498071 [=====] - 2s 0us/step

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
'dog', 'frog', 'horse', 'ship', 'truck']
```

2s completed at 21:49

practical 3(b,c).ipynb

```
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    # The CIFAR labels happen to be arrays,
    # which is why you need the extra index
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()
```

2s completed at 21:49

practical 3(b,c).ipynb

The screenshot shows a Jupyter Notebook interface with three tabs visible at the top: 'My Drive - Go...', 'Common Doc...', and 'practical 3(b,c).ipynb'. The notebook content includes Python code for building a neural network and a 'model.summary()' command.

```

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

{x} [ ] model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 6, 6, 64)	0	
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928

✓ 2s completed at 21:49

The second screenshot shows the same Jupyter Notebook environment. The code and model summary output are identical to the first one.

```

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

{x} [ ] model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 6, 6, 64)	0	
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
flatten (Flatten)	(None, 1024)	0
dense_3 (Dense)	(None, 64)	65600

✓ 2s completed at 21:49

The third screenshot shows the same Jupyter Notebook environment. The code and model summary output are identical to the previous ones.

```

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

{x} [ ] model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 6, 6, 64)	0	
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
flatten (Flatten)	(None, 1024)	0
dense_3 (Dense)	(None, 64)	65600

✓ 2s completed at 21:49

The screenshot shows the Google Colab interface with the notebook titled "practical3.ipynb". The code cell [7] contains the following Python code:

```
[7]: Epoch 150/150  
21s 77/77 [=====] - 0s 2ms/step - loss: 0.4835 - accuracy: 0.7617  
<keras.callbacks.History at 0x7f79d5022190>
```

The code cell [8] contains:

```
[8]: _,accuracy=model.evaluate(X,Y)
```

The code cell [9] contains:

```
[9]: print('accuracy of model is',(accuracy*100) )
```

The output for cell [9] is:

```
accuracy of model is 76.5625
```

The code cell [11] contains:

```
[11]: predictions = model.predict(X)
```

```
#Compile and train the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                     validation_data=(test_images, test_labels))

Epoch 1/10
1563/1563 [=====] - 72s 46ms/step - loss: 1.4818 - accuracy: 0.4596 - val_loss: 1.2175 - val_accuracy: 0.5641
Epoch 2/10
1563/1563 [=====] - 72s 46ms/step - loss: 1.1231 - accuracy: 0.6012 - val_loss: 1.0866 - val_accuracy: 0.6181
Epoch 3/10
1563/1563 [=====] - 72s 46ms/step - loss: 0.9719 - accuracy: 0.6598 - val_loss: 1.0174 - val_accuracy: 0.6382
Epoch 4/10
1563/1563 [=====] - 71s 46ms/step - loss: 0.8801 - accuracy: 0.6920 - val_loss: 0.8800 - val_accuracy: 0.6970
Epoch 5/10
1563/1563 [=====] - 69s 44ms/step - loss: 0.7982 - accuracy: 0.7238 - val_loss: 0.9590 - val_accuracy: 0.6747
Epoch 6/10
```

My Dr x Comm x Comm x lonos x pima... x Practi x +

colab.research.google.com/drive/1QjczT9FU2RQv8GGzx1O6SUWQX12Mjxn#scrollTo=abW8Oy2Pf0kB

Apps Gmail New Tab Classes Introduction to Am...

practical 3(b,c).ipynb

File Edit View Insert Runtime Tools Help Cannot save changes

RAM Disk Editing

1563/1563 [=====] - 70s 45ms/step - loss: 0.7319 - accuracy: 0.7435 - val_loss: 0.9055 - val_accuracy: 0.6870

Epoch 8/10

1563/1563 [=====] - 70s 45ms/step - loss: 0.6830 - accuracy: 0.7595 - val_loss: 0.8536 - val_accuracy: 0.7163

Epoch 9/10

1563/1563 [=====] - 70s 45ms/step - loss: 0.6397 - accuracy: 0.7743 - val_loss: 0.9195 - val_accuracy: 0.6927

Epoch 10/10

1563/1563 [=====] - 70s 45ms/step - loss: 0.6054 - accuracy: 0.7862 - val_loss: 0.8725 - val_accuracy: 0.7092

#Evaluate the model

```
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```

313/313 - 4s - loss: 0.8725 - accuracy: 0.7092 - 4s/epoch - 12ms/step

10

✓ 4s completed at 22:06

practical_5(a,b,c).py lonospear.csv pima-indians-diab...csv

Type here to search EN 30°C AQI 133 22:19 26-11-2021

colab.research.google.com/drive/1QjczT9FU2RQv8GGzx1O6SUWQX12Mjxn#scrollTo=abW8Oy2Pf0kB

Apps Gmail New Tab Classes Introduction to Am...

practical 3(b,c).ipynb

File Edit View Insert Runtime Tools Help Cannot save changes

RAM Disk Editing

plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)

313/313 - 4s - loss: 0.8725 - accuracy: 0.7092 - 4s/epoch - 12ms/step

{x}

10

0.9

0.8

0.7

0.6

0.5

Accuracy

Epoch

✓ 4s completed at 22:06

practical_5(a,b,c).py lonospear.csv pima-indians-diab...csv

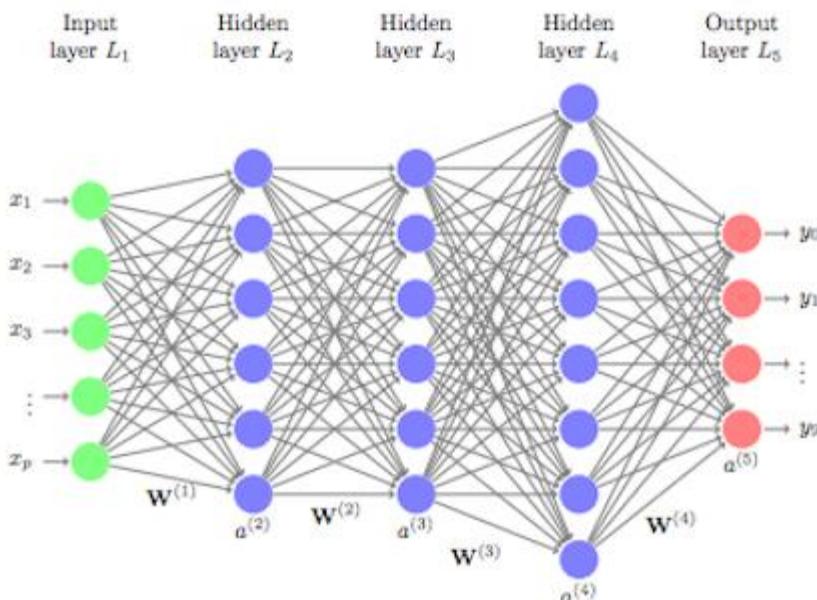
Type here to search EN 30°C Haze 22:20 26-11-2021

Practical No :- 4

**Aim: Using deep feed forward network with two hidden layers
for performing classification and predicting the class.**

https://colab.research.google.com/drive/1lrlIB3GrdQ9L9KVSm_TDfXNz8Bh4_dx

Machine learning algorithms typically search for the optimal representation of data using some feedback signal (aka objective/loss function). However, most machine learning algorithms only have the ability to use one or two layers of data transformation to learn the output representation. As data sets continue to grow in the dimensions of the feature space, finding the optimal output representation with a *shallow* model is not always possible. Deep learning provides a multi-layer approach to learn data representations, typically performed with a *multi-layer neural network*. Like other machine learning algorithms, deep neural networks (DNN) perform learning by mapping features to targets through a process of simple data transformations and feedback signals; however, DNNs place an emphasis on learning successive layers of meaningful representations. Although an intimidating subject, the overarching concept is rather simple and has proven highly successful in predicting a wide range of problems (i.e. image classification, speech recognition, autonomous driving).



Code:-

```
from keras.models import Sequential
from keras.layers import Dense
from sklearn.datasets import make_blobs
from sklearn.preprocessing import MinMaxScaler
import numpy as np

X,Y=make_blobs(n_samples=100,centers=2,n_features=2,random_state=1)
```

```

scalar=MinMaxScaler()
scalar.fit(X)
X=scalar.transform(X)

model=Sequential()
model.add(Dense(4,input_dim=2,activation='relu'))
model.add(Dense(4,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

model.compile(loss='binary_crossentropy',optimizer='adam')
model.fit(X,Y,epochs=500)

Xnew,Yreal=make_blobs(n_samples=5,centers=2,n_features=2,random_state=1)

Xnew=scalar.transform(Xnew)
Ynew=model.predict(Xnew)
Ynew=np.round(Ynew).astype(int)

for i in range(len(Xnew)):
    print("X=%s,Predicted=%d,Desired=%s"%(Xnew[i],Ynew[i],Yreal[i]))

```

The screenshot shows a Google Colab interface with multiple tabs open. The active tab is titled "Practical No:4.ipynb". The code cell contains the Python script provided above. The output pane shows the code being executed and completed successfully at 22:25. The status bar at the bottom indicates "16s completed at 22:25".

```

#Aim: Using deep feed forward network with two hidden layers for performing classification and predicting the class.

from keras.models import Sequential
from keras.layers import Dense
from sklearn.datasets import make_blobs
from sklearn.preprocessing import MinMaxScaler
import numpy as np

X,Y=make_blobs(n_samples=100,centers=2,n_features=2,random_state=1)
scalar=MinMaxScaler()
scalar.fit(X)
X=scalar.transform(X)

model=Sequential()
model.add(Dense(4,input_dim=2,activation='relu'))
model.add(Dense(4,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

model.compile(loss='binary_crossentropy',optimizer='adam')
model.fit(X,Y,epochs=500)

```

Output:-

```

Practical No:4.ipynb
File Edit View Insert Runtime Tools Help Cannot save changes
+ Code + Text Copy to Drive
epoch 491/500
4/4 [=====] - 0s 4ms/step - loss: 0.1024
Epoch 492/500
4/4 [=====] - 0s 4ms/step - loss: 0.1021
Epoch 493/500
4/4 [=====] - 0s 4ms/step - loss: 0.1019
Epoch 494/500
4/4 [=====] - 0s 5ms/step - loss: 0.1016
Epoch 495/500
4/4 [=====] - 0s 7ms/step - loss: 0.1013
Epoch 496/500
4/4 [=====] - 0s 6ms/step - loss: 0.1011
Epoch 497/500
4/4 [=====] - 0s 6ms/step - loss: 0.1008
Epoch 498/500
4/4 [=====] - 0s 7ms/step - loss: 0.1006
Epoch 499/500
4/4 [=====] - 0s 5ms/step - loss: 0.1003
Epoch 500/500
4/4 [=====] - 0s 4ms/step - loss: 0.1000
X=[0.26814469 0.0290577 ],Predicted=1,Desired=1
X=[0.17582555 0.16948267],Predicted=1,Desired=1
X=[0.89337759 0.65864154],Predicted=0,Desired=0
X=[0.96440204 0.77809405],Predicted=0,Desired=0
X=[0.78082614 0.75391697],Predicted=0,Desired=0

```

b) Aim: Using a deep field forward network with two hidden layers for performing classification and predicting the probability of class.

Code:-

```

from keras.models import Sequential
from keras.layers import Dense
from sklearn.datasets import make_blobs
from sklearn.preprocessing import MinMaxScaler

X,Y=make_blobs(n_samples=100,centers=2,n_features=2,random_state=1)
scalar=MinMaxScaler()
scalar.fit(X)
X=scalar.transform(X)

model=Sequential()
model.add(Dense(4,input_dim=2,activation='relu'))
model.add(Dense(4,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

model.compile(loss='binary_crossentropy',optimizer='adam')
model.fit(X,Y,epochs=500)

Xnew,Yreal=make_blobs(n_samples=3,centers=2,n_features=2,random_state=1)
Xnew=scalar.transform(Xnew)
Yclass=model.predict(Xnew)

```

```

Ynew=model.predict(Xnew)
for i in range(len(Xnew)):
    print("X=%s,Predicted_probability=%s,Predicted_class=%s"%(Xnew[i],Ynew[i],Yclass[i]))

```

Output:-:

```

#Practical 4b
#b) Aim: Using a deep field forward network with two hidden layers for performing classification and predicting the probability of class.

from keras.models import Sequential
from keras.layers import Dense
from sklearn.datasets import make_blobs
from sklearn.preprocessing import MinMaxScaler

X,Y=make_blobs(n_samples=100,centers=2,n_features=2,random_state=1)
scalar=MinMaxScaler()
scalar.fit(X)
X=scalar.transform(X)

model=Sequential()
model.add(Dense(4,input_dim=2,activation='relu'))
model.add(Dense(4,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

model.compile(loss='binary_crossentropy',optimizer='adam')
model.fit(X,Y,epochss500)

Xnew,Yreal=make_blobs(n_samples=3,centers=2,n_features=2,random_state=1)
Xnew=scalar.transform(Xnew)

```

```

4/4 [=====] - 0s 5ms/step - loss: 0.0051
Epoch 491/500
4/4 [=====] - 0s 5ms/step - loss: 0.0051
Epoch 492/500
4/4 [=====] - 0s 5ms/step - loss: 0.0051
Epoch 493/500
4/4 [=====] - 0s 5ms/step - loss: 0.0050
Epoch 494/500
4/4 [=====] - 0s 5ms/step - loss: 0.0050
Epoch 495/500
4/4 [=====] - 0s 4ms/step - loss: 0.0050
Epoch 496/500
4/4 [=====] - 0s 5ms/step - loss: 0.0050
Epoch 497/500
4/4 [=====] - 0s 6ms/step - loss: 0.0049
Epoch 498/500
4/4 [=====] - 0s 4ms/step - loss: 0.0049
Epoch 499/500
4/4 [=====] - 0s 4ms/step - loss: 0.0049
Epoch 500/500
4/4 [=====] - 0s 4ms/step - loss: 0.0048
X=[0.89337759 0.65864154],Predicted_probability=[0.00317386],Predicted_class=[0.00317386]
X=[0.29097707 0.12978982],Predicted_probability=[0.99100596],Predicted_class=[0.99100596]
X=[0.78082614 0.75391697],Predicted_probability=[0.00572661],Predicted_class=[0.00572661]

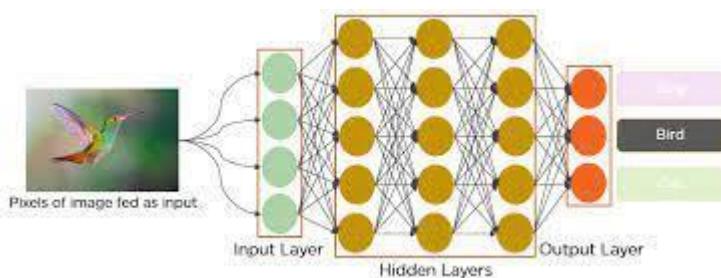
```

Practical No :- 5

Aim: cnn of cifer10 images

In the past few decades, Deep Learning has proved to be a very powerful tool because of its ability to handle large amounts of data. The interest to use hidden layers has surpassed traditional techniques, especially in pattern recognition. One of the most popular deep neural networks is Convolutional Neural Networks.

<https://colab.research.google.com/drive/1C1hqqu2qpSzuGPKUohReCgIrmRcY-kH#scrollTo=sjoN8Bjo4P-2>



Code:-

```
import tensorflow as tf
```

```

from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt

(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    # The CIFAR labels happen to be arrays,
    # which is why you need the extra index
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))

model.summary()
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=10,
                     validation_data=(test_images, test_labels))
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')

test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)

```

Output:-

The screenshot shows two consecutive screenshots of a Google Colab notebook titled "Practical 5(a,b,c)".

Screenshot 1: The code cell contains the following Python code to load CIFAR-10 images:

```
[1]: import tensorflow as tf  
from tensorflow.keras import datasets, layers, models  
import matplotlib.pyplot as plt  
  
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()  
  
# Normalize pixel values to be between 0 and 1  
train_images, test_images = train_images / 255.0, test_images / 255.0
```

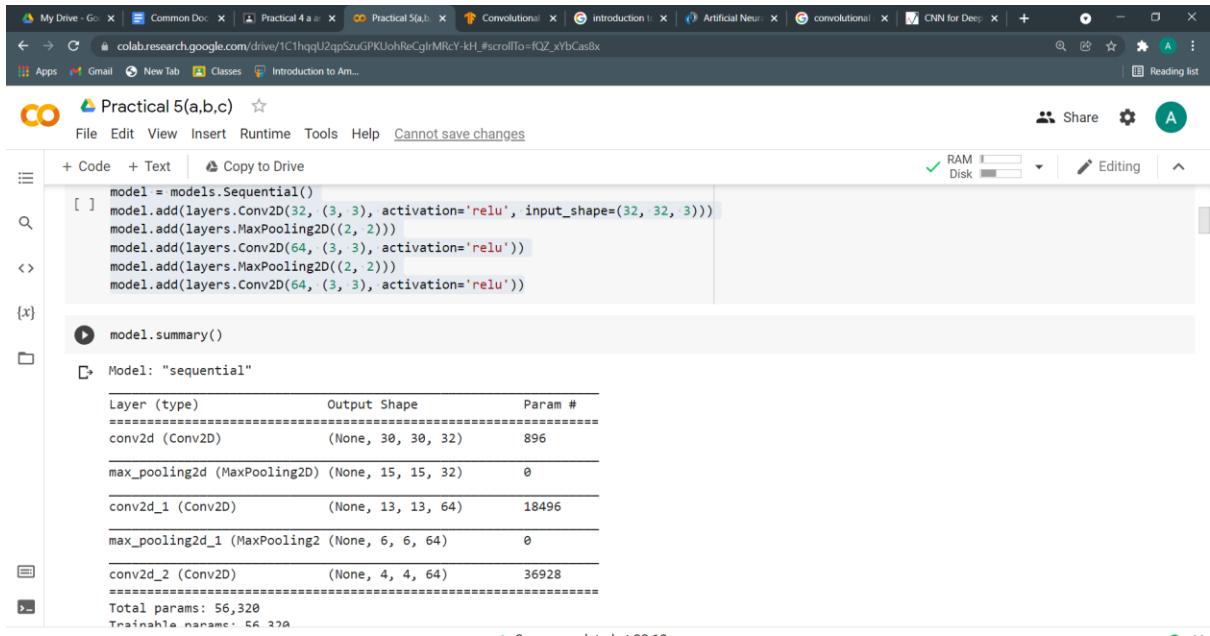
Output: Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 4s 0us/step
170500096/170498071 [=====] - 4s 0us/step
✓ 2s completed at 22:12

Screenshot 2: The code cell contains the following Python code to display the first image in the training set:

```
plt.imshow(train_images[0])  
plt.show()
```

Output: A 4x5 grid of 20 CIFAR-10 images with their corresponding labels below them:

frog	truck	truck	deer	automobile
automobile	bird	horse	ship	cat
deer	horse	horse	bird	truck
building	house	bird	apple	truck



```

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

```

{x}

```

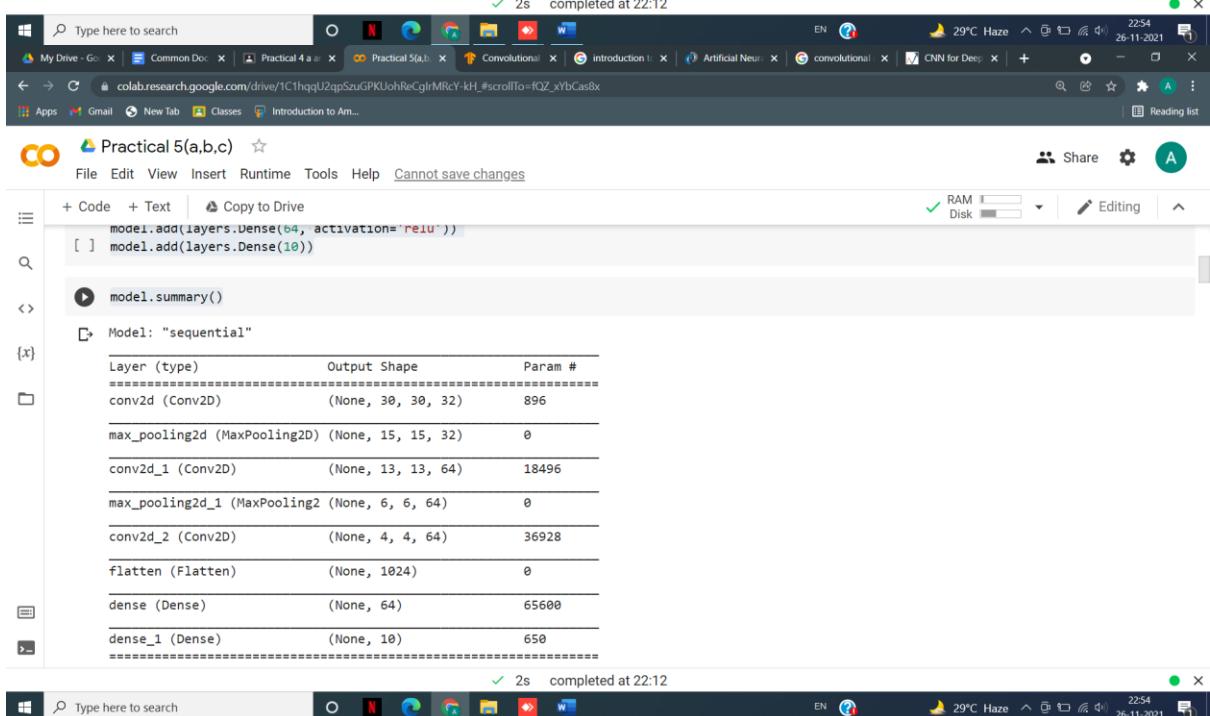
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
Total params:	56,320	
Trainable params:	56,220	

✓ 2s completed at 22:12



```

model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))

```

{x}

```

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 64)	65600
dense_1 (Dense)	(None, 10)	650

✓ 2s completed at 22:12



b. Image classification

Code:-

```

import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
flower_photo/

daisy/
dandelion/
roses/
sunflowers/
tulips/
import pathlib
dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
"
data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, untar=True)
data_dir = pathlib.Path(data_dir)

```

```

image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)
roses = list(data_dir.glob('roses/*'))
PIL.Image.open(str(roses[0]))
PIL.Image.open(str(roses[1]))
tulips = list(data_dir.glob('tulips/*'))
PIL.Image.open(str(tulips[0]))
PIL.Image.open(str(tulips[1]))
batch_size = 32
img_height = 180
img_width = 180
train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
class_names = train_ds.class_names
print(class_names)
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
    for image_batch, labels_batch in train_ds:
        print(image_batch.shape)
        print(labels_batch.shape)
        break
AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
normalization_layer = layers.Rescaling(1./255)
normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))

```

```

image_batch, labels_batch = next(iter(normalized_ds))
first_image = image_batch[0]
# Notice the pixel values are now in '[0,1].
print(np.min(first_image), np.max(first_image))
num_classes = 5

model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
epochs=10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
)
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
data_augmentation = keras.Sequential(

```

```

[
    layers.RandomFlip("horizontal",
        input_shape=(img_height,
                    img_width,
                    3)),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.1),
]
)
plt.figure(figsize=(10, 10))
for images, _ in train_ds.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")
model = Sequential([
    data_augmentation,
    layers.Rescaling(1./255),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
epochs = 10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

```

```

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

sunflower_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg"
sunflower_path = tf.keras.utils.get_file('Red_sunflower', origin=sunflower_url)

img = tf.keras.utils.load_img(
    sunflower_path, target_size=(img_height, img_width)
)
img_array = tf.keras.utils.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)

```

Output-:

My Drive | Common D | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Deep Learning | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglMRcY-kH.#scrollTo=qpXSFtke25Q9

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

Image classification

<https://www.tensorflow.org/tutorials/images/classification>

```

import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential

```

flower_photo/

- daisy/
- dandelion/

2s completed at 22:12

Practical_5(a,b,c).ipynb | Practical_No_4.ipynb | Show all

Type here to search

EN 29°C Haze 23:07 26-11-2021

My Drive | Common D | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Deep Learning | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglMRcY-kH.#scrollTo=qpXSFtke25Q9

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

```

dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, untar=True)
data_dir = pathlib.Path(data_dir)

```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz
228818944/228813984 [=====] - 2s 0us/step
228827136/228813984 [=====] - 2s 0us/step

{x}

```

image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)

```

3670

```

[ ] roses = list(data_dir.glob('roses/*'))
PIL.Image.open(str(roses[0]))

```



2s completed at 22:12

Practical_5(a,b,c).ipynb | Practical_No_4.ipynb | Show all

Type here to search

EN 29°C Haze 23:08 26-11-2021

My Drive | Common | Practical 4 | Practical Note | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D... | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhohReCglrMRcY-kH.#scrollTo=qpXSFkze25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

PIL.Image.open(str(roses[1]))

2s completed at 22:12

Type here to search

EN 29°C Haze 23:08 26-11-2021

My Drive | Common | Practical 4 | Practical Note | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D... | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhohReCglrMRcY-kH.#scrollTo=qpXSFkze25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

tulips = list(data_dir.glob('tulips/*'))
PIL.Image.open(str(tulips[0]))

2s completed at 22:12

Type here to search

EN 29°C Haze 23:08 26-11-2021

My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Dummies | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhohReCglrMRcY-kH.#scrollTo=qpXSF1ke25Q9

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ]
```



```
[ ] batch_size = 32
img_height = 180
img_width = 180

train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

2s completed at 22:12

Type here to search

EN 29°C Haze 23:08 26-11-2021

My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Dummies | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhohReCglrMRcY-kH.#scrollTo=qpXSF1ke25Q9

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] class_names = train_ds.class_names
[ ] print(class_names)

['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']

import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```

2s completed at 22:12

roses dandelion tulips

Type here to search

EN 29°C Haze 23:08 26-11-2021

My Drive | Common | Practical 4 | Practical N | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhReCglrMRcY-kH.#scrollTo=qpXSFkze25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

dandelion roses tulips

```
[ ] for image_batch, labels_batch in train_ds:
    print(image_batch.shape)
    print(labels_batch.shape)
    break

(32, 180, 180, 3)
(32,)
```

Configure the dataset for performance

Type here to search 2s completed at 22:12

EN 29°C Haze 23:08 26-11-2021

My Drive | Common | Practical 4 | Practical N | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUhReCglrMRcY-kH.#scrollTo=qpXSFkze25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

RAM Disk Editing

Configure the dataset for performance

```
[ ] AUTOTUNE = tf.data.AUTOTUNE

{ } train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
{ } val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

[ ] normalization_layer = layers.Rescaling(1./255)

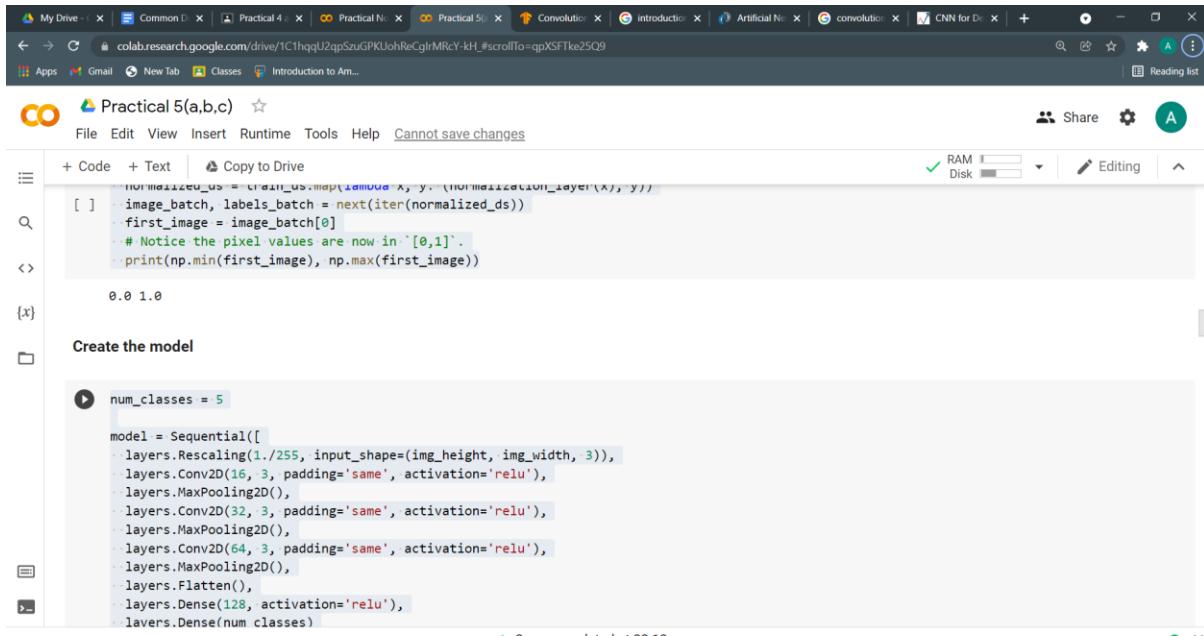
[ ] normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
[ ] image_batch, labels_batch = next(iter(normalized_ds))
[ ] first_image = image_batch[0]
[ ] # Notice the pixel values are now in '[0,1].
[ ] print(np.min(first_image), np.max(first_image))

0.0 1.0
```

Create the model

Type here to search 2s completed at 22:12

EN 29°C Haze 23:08 26-11-2021



```

My Drive | Common D... | Practical 4... | Practical N... | Practical 5... | Convolution... | introduction... | Artificial Ne... | convolution... | CNN for D... | + | - | X | Reading list
← → 🔒 colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglrMRcY-kH.#scrollTo=qpXSF1ke25Q9
Apps Gmail New Tab Classes Introduction to Am...
CO Practical 5(a,b,c) ★
File Edit View Insert Runtime Tools Help Cannot save changes
+ Code + Text ⚡ Copy to Drive
[ ] · normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
· image_batch, labels_batch = next(iter(normalized_ds))
· first_image = image_batch[0]
· # Notice the pixel values are now in '[0,1]'.
· print(np.min(first_image), np.max(first_image))

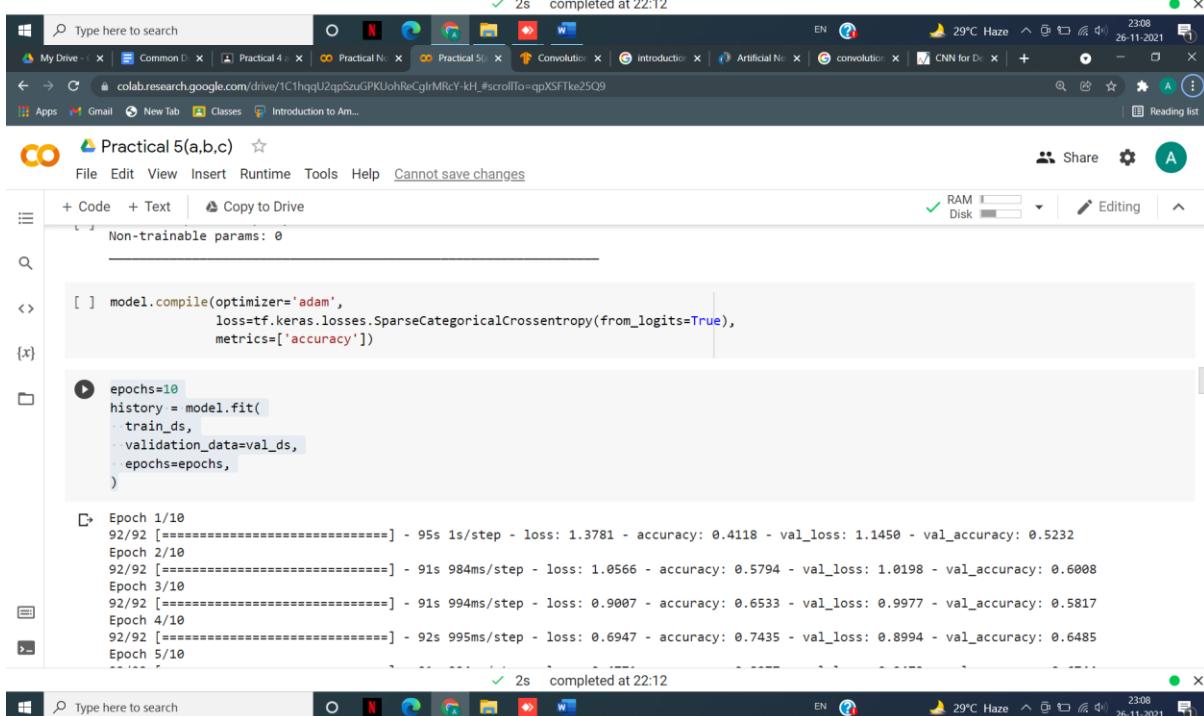
0.0 1.0
{x}

Create the model

▶ num_classes = 5

model = Sequential([
· layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
· layers.Conv2D(16, 3, padding='same', activation='relu'),
· layers.MaxPooling2D(),
· layers.Conv2D(32, 3, padding='same', activation='relu'),
· layers.MaxPooling2D(),
· layers.Conv2D(64, 3, padding='same', activation='relu'),
· layers.MaxPooling2D(),
· layers.Flatten(),
· layers.Dense(128, activation='relu'),
· layers.Dense(num_classes)
]
2s completed at 22:12

```



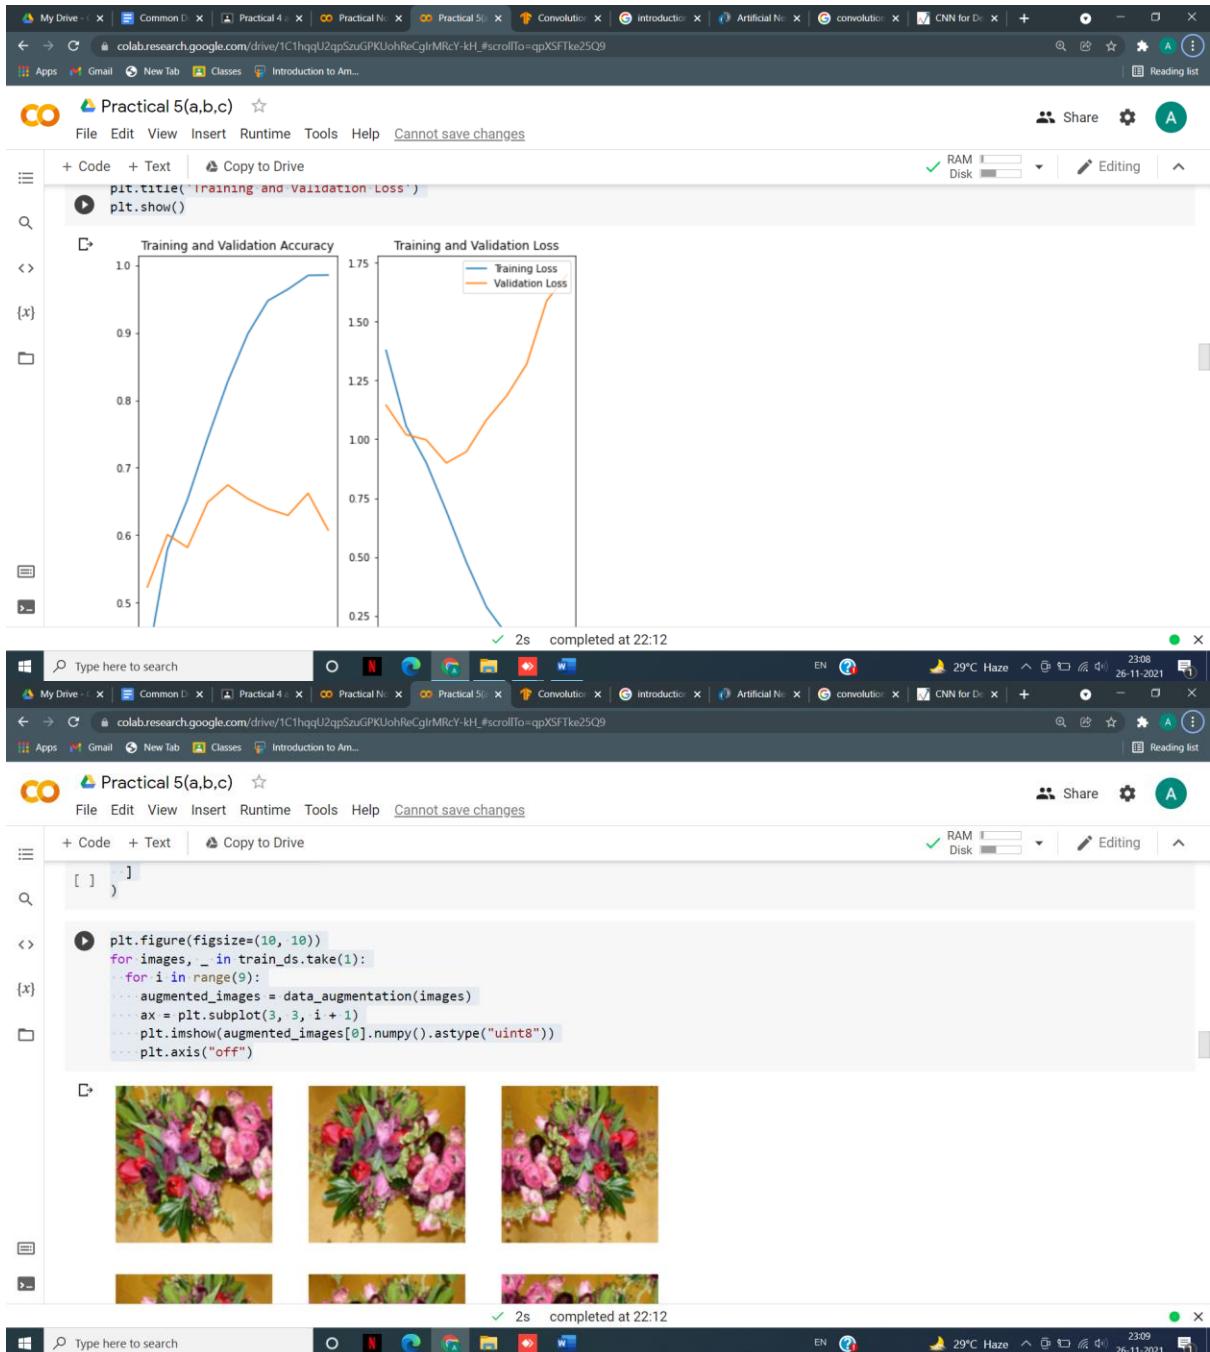
```

Type here to search O Netflix Google Photos W
My Drive | Common D... | Practical 4... | Practical N... | Practical 5... | Convolution... | introduction... | Artificial Ne... | convolution... | CNN for D... | + | - | X | Reading list
← → 🔒 colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglrMRcY-kH.#scrollTo=qpXSF1ke25Q9
Apps Gmail New Tab Classes Introduction to Am...
CO Practical 5(a,b,c) ★
File Edit View Insert Runtime Tools Help Cannot save changes
+ Code + Text ⚡ Copy to Drive
Non-trainable params: 0
[ ] model.compile(optimizer='adam',
· loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
· metrics=['accuracy'])

{x}

▶ epochs=10
history = model.fit(
· train_ds,
· validation_data=val_ds,
· epochs=epochs,
)
Epoch 1/10
92/92 [=====] - 95s 1s/step - loss: 1.3781 - accuracy: 0.4118 - val_loss: 1.1450 - val_accuracy: 0.5232
Epoch 2/10
92/92 [=====] - 91s 984ms/step - loss: 1.0566 - accuracy: 0.5794 - val_loss: 1.0198 - val_accuracy: 0.6008
Epoch 3/10
92/92 [=====] - 91s 994ms/step - loss: 0.9007 - accuracy: 0.6533 - val_loss: 0.9977 - val_accuracy: 0.5817
Epoch 4/10
92/92 [=====] - 92s 995ms/step - loss: 0.6947 - accuracy: 0.7435 - val_loss: 0.8994 - val_accuracy: 0.6485
Epoch 5/10
2s completed at 22:12

```



My Drive | Common D | Practical 4 | Practical Note | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglrMRcY-kH.#scrollTo=qpXSFtke25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] layers.MaxPooling2D(),
[ ] layers.Conv2D(32, 3, padding='same', activation='relu'),
[ ] layers.MaxPooling2D(),
[ ] layers.Conv2D(64, 3, padding='same', activation='relu'),
[ ] layers.MaxPooling2D(),
[ ] layers.Dropout(0.2),
[ ] layers.Flatten(),
[ ] layers.Dense(128, activation='relu'),
[ ] layers.Dense(num_classes)
[])

[ ] model.compile(optimizer='adam',
[ ] loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
[ ] metrics=['accuracy'])

▶ epochs = 10
history = model.fit(
[ ] train_ds,
[ ] validation_data=val_ds,
[ ] epochs=epochs
)
```

2s completed at 22:12

Type here to search

EN 29°C Haze 23:09 26-11-2021

My Drive | Common D | Practical 4 | Practical Note | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCglrMRcY-kH.#scrollTo=qpXSFtke25Q9

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
▶ plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```

Training and Validation Accuracy

Training and Validation Loss

2s completed at 22:12

Type here to search

EN 29°C Haze 23:09 26-11-2021

The screenshot shows a Google Colab notebook titled "Practical 5(a,b,c)". The code cell contains Python code for loading a sunflower image, processing it with TensorFlow, and printing the prediction results. A progress bar at the top indicates "Validation Accuracy" from 0 to 10, with a value of 0.5 highlighted. Below the code, a terminal window shows the command to download the image and its progress.

```

sunflower_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg"
sunflower_path = tf.keras.utils.get_file('Red_sunflower', origin=sunflower_url)

img = tf.keras.utils.load_img(
    sunflower_path, target_size=(img_height, img_width)
)
img_array = tf.keras.utils.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)

```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg
122880/117948 [=====] - 0s 0us/step
131072/117948 [=====] - 0s 0us/step ✓ 2s completed at 22:12

c.Data Augmentation

Code:-

```

import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds

from tensorflow.keras import layers
(train_ds, val_ds, test_ds), metadata = tfds.load(
    'tf_flowers',
    split=['train[:80%]', 'train[80%:90%]', 'train[90%:]'],
    with_info=True,
    as_supervised=True,
)
num_classes = metadata.features['label'].num_classes
print(num_classes)
get_label_name = metadata.features['label'].int2str

image, label = next(iter(train_ds))
_=plt.imshow(image)
_=plt.title(get_label_name(label))
IMG_SIZE = 180

```

```

resize_and_rescale = tf.keras.Sequential([
    layers.Resizing(IMG_SIZE, IMG_SIZE),
    layers.Rescaling(1./255)
])
result = resize_and_rescale(image)
_ = plt.imshow(result)
print("Min and max pixel values:", result.numpy().min(), result.numpy().max())
data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
])
# Add the image to a batch.
image = tf.expand_dims(image, 0)
plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_image = data_augmentation(image)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_image[0])
    plt.axis("off")
model = tf.keras.Sequential([
    # Add the preprocessing layers you created earlier.
    resize_and_rescale,
    data_augmentation,
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    # Rest of your model.
])
aug_ds = train_ds.map(
    lambda x, y: (resize_and_rescale(x, training=True), y))
batch_size = 32
AUTOTUNE = tf.data.AUTOTUNE

def prepare(ds, shuffle=False, augment=False):
    # Resize and rescale all datasets.
    ds = ds.map(lambda x, y: (resize_and_rescale(x), y),
               num_parallel_calls=AUTOTUNE)

    if shuffle:
        ds = ds.shuffle(1000)

    # Batch all datasets.
    ds = ds.batch(batch_size)

    # Use data augmentation only on the training set.
    if augment:
        ds = ds.map(lambda x, y: (data_augmentation(x, training=True), y),

```

```

    num_parallel_calls=AUTOTUNE)

# Use buffered prefetching on all datasets.
return ds.prefetch(buffer_size=AUTOTUNE)
train_ds = prepare(train_ds, shuffle=True, augment=True)
val_ds = prepare(val_ds)
test_ds = prepare(test_ds)
model = tf.keras.Sequential([
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
epochs=5
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
def random_invert_img(x, p=0.5):
    if tf.random.uniform([]) < p:
        x = (255-x)
    else:
        x
    return x
def random_invert(factor=0.5):
    return layers.Lambda(lambda x: random_invert_img(x, factor))

random_invert = random_invert()
plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_image = random_invert(image)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_image[0].numpy().astype("uint8"))
    plt.axis("off")
class RandomInvert(layers.Layer):
    def __init__(self, factor=0.5, **kwargs):
        super().__init__(**kwargs)

```

```

self.factor = factor

def call(self, x):
    return random_invert_img(x)
_ = plt.imshow(RandomInvert()(image)[0])
(train_ds, val_ds, test_ds), metadata = tfds.load(
    'tf_flowers',
    split=['train[:80%]', 'train[80%:90%]', 'train[90%:]'],
    with_info=True,
    as_supervised=True,
)
image, label = next(iter(train_ds))
_ = plt.imshow(image)
_ = plt.title(get_label_name(label))
def visualize(original, augmented):
    fig = plt.figure()
    plt.subplot(1,2,1)
    plt.title('Original image')
    plt.imshow(original)

    plt.subplot(1,2,2)
    plt.title('Augmented image')
    plt.imshow(augmented)
    # flipped = tf.image.flip_left_right(image)

    visualize(image, flipped)
grayscaled = tf.image.rgb_to_grayscale(image)
visualize(image, tf.squeeze(grayscaled))
_ = plt.colorbar()
saturated = tf.image.adjust_saturation(image, 3)
visualize(image, saturated)
Flip an image

```

Flip an image either vertically or horizontally with `tf.image.flip_left_right`:

```

for i in range(3):
    seed = (i, 0) # tuple of size (2,)
    stateless_random_contrast = tf.image.stateless_random_contrast(
        image, lower=0.1, upper=0.9, seed=seed)
    visualize(image, stateless_random_contrast)
#### Randomly crop an image

```

##Randomly crop `image` using `tf.image.stateless_random_crop` by providing target `size` and `seed`. The portion that gets cropped out of `image` is at a randomly chosen offset and is associated with the given `seed`.

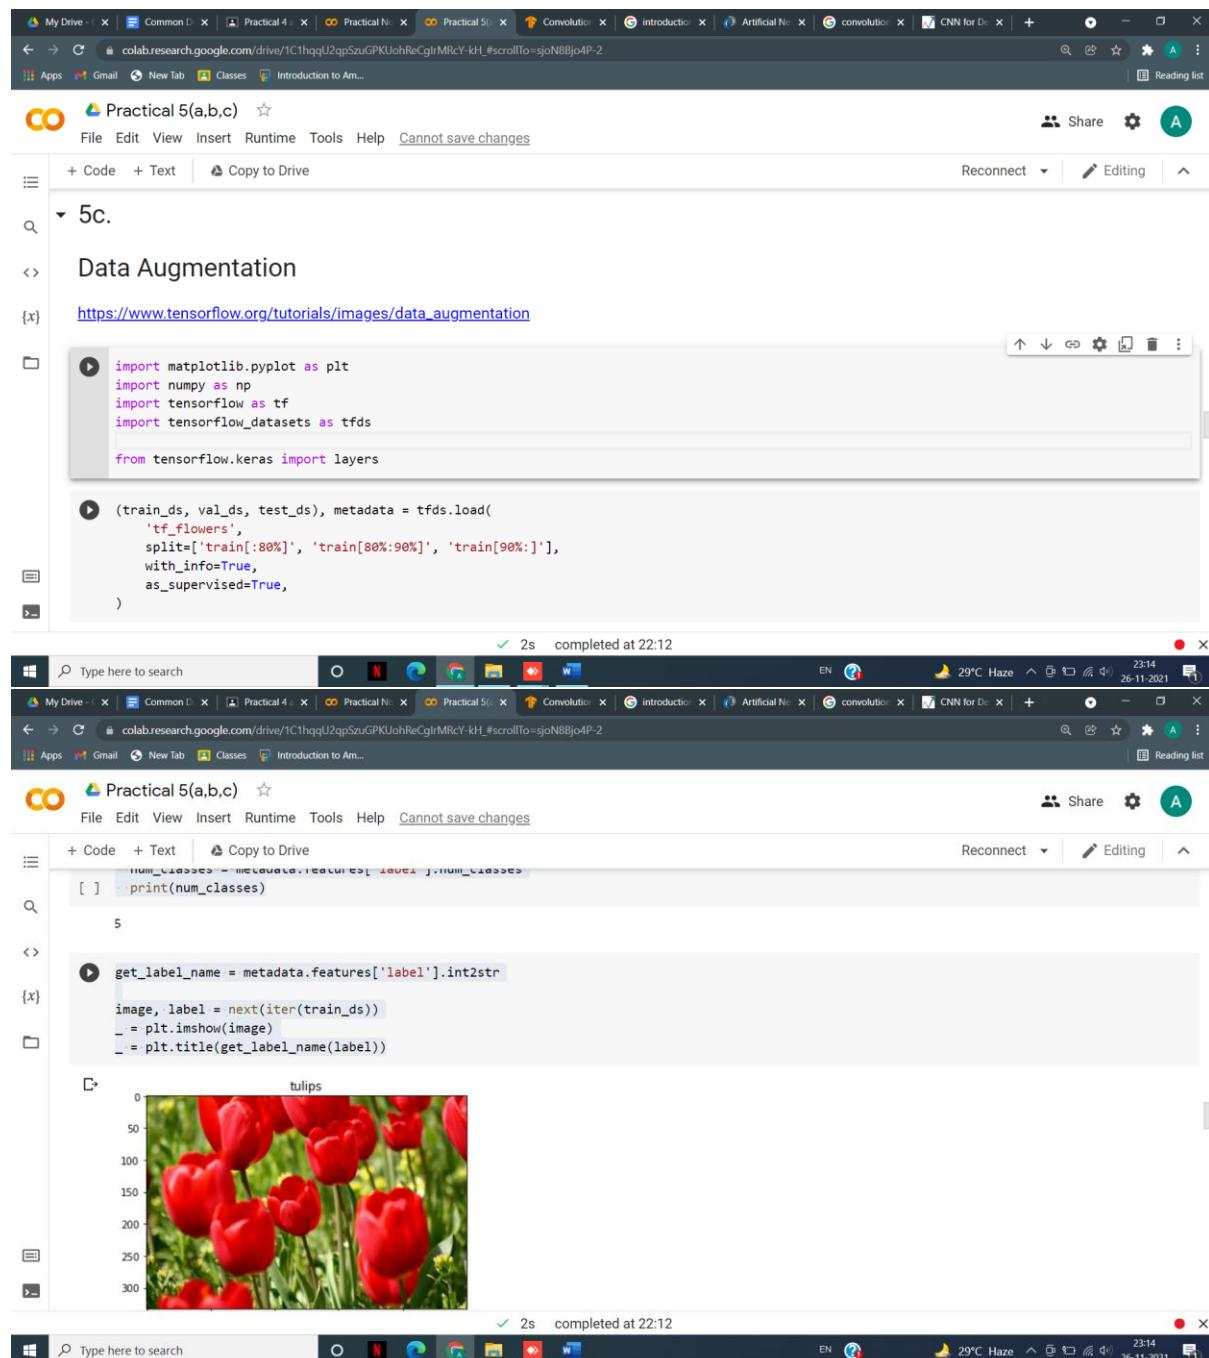
```

for i in range(3):
    seed = (i, 0) # tuple of size (2,)
    stateless_random_crop = tf.image.stateless_random_crop(
        image, size=[210, 300, 3], seed=seed)

```

```
visualize(image, stateless_random_crop)
```

Output:-



The screenshot shows two consecutive executions of code in a Google Colab notebook titled "Practical 5(a,b,c)".

Execution 1:

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds

from tensorflow.keras import layers
```

Execution 2:

```
[ ] num_classes = metadata.features['label'].num_classes
[ ] print(num_classes)

5

[ ] get_label_name = metadata.features['label'].int2str
[ ] image, label = next(iter(train_ds))
[ ] _ = plt.imshow(image)
[ ] _ = plt.title(get_label_name(label))

[ ] tulips
```

A thumbnail image of red tulips is displayed below the second cell's output.

My Drive | Common D | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for D | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
IMG_SIZE = 180
[ ] resize_and_rescale = tf.keras.Sequential([
    layers.Resizing(IMG_SIZE, IMG_SIZE),
    layers.Rescaling(1./255)
])
{x}
[ ] result = resize_and_rescale(image)
_ = plt.imshow(result)

[ ] plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_image = data_augmentation(image)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_image[0])
    plt.axis("off")
```

Reconnect Editing

2s completed at 22:12

Type here to search

EN 29°C Haze 23:14 26-11-2021

My Drive | Common D | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for D | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
Min and max pixel values: 0.0 1.0
[ ] data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
])
{x}
[ ] # Add the image to a batch.
image = tf.expand_dims(image, 0)

[ ] plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_image = data_augmentation(image)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_image[0])
    plt.axis("off")
```

Reconnect Editing

2s completed at 22:12

Type here to search

EN 29°C Haze 23:14 26-11-2021

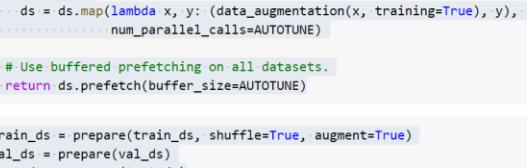


```

model = tf.keras.Sequential([
    # Add the preprocessing layers you created earlier.
    resize_and_rescale,
    data_augmentation,
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    # Rest of your model.
])

```

2s completed at 22:12



```

ds = ds.map(lambda x, y: (data_augmentation(x, training=True), y),
            num_parallel_calls=AUTOTUNE)

# Use buffered prefetching on all datasets.
return ds.prefetch(buffer_size=AUTOTUNE)

```

2s completed at 22:12



```

train_ds = prepare(train_ds, shuffle=True, augment=True)
val_ds = prepare(val_ds)
test_ds = prepare(test_ds)

model = tf.keras.Sequential([
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])

```

2s completed at 22:12

My Drive | Common D | Practical 4 | Practical N | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hpqlJ2qpSzGPKUohReCgirMRcY-kH_.#scrollTo=sjoN8Bjo4P-2

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] return x

def random_invert(factor=0.5):
    return layers.Lambda(lambda x: random_invert_img(x, factor))

random_invert = random_invert()

plt.figure(figsize=(10, 10))
for i in range(9):
    augmented_image = random_invert(image)
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(augmented_image[0].numpy().astype("uint8"))
    plt.axis("off")
```

2s completed at 22:12

Type here to search

EN 29°C Haze 23:14 26-11-2021

My Drive | Common D | Practical 4 | Practical N | Practical 5 | Convolution | introduction | Artificial Ne | convolution | CNN for D | +

colab.research.google.com/drive/1C1hpqlJ2qpSzGPKUohReCgirMRcY-kH_.#scrollTo=sjoN8Bjo4P-2

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] def __init__(self, factor=0.5, **kwargs):
    super().__init__(**kwargs)
    self.factor = factor

    def call(self, x):
        return random_invert_img(x)

{x}
plt.imshow(RandomInvert()(image)[0])
```

2s completed at 22:12

Type here to search

EN 29°C Haze 23:14 26-11-2021

My Drive | Common | Practical 4 | Practical Note | Practical 5 | Convolution | introduction | Artificial Neuron | convolution | CNN for D... | +

colab.research.google.com/drive/1C1hgqlU2qpSzUGPKUhReCglrMRcY-kH.#scrollTo=sjoN8Bjjo4P-2

Apps Gmail New Tab Classes Introduction to Am...

Reading list

Practical 5(a,b,c) ☆

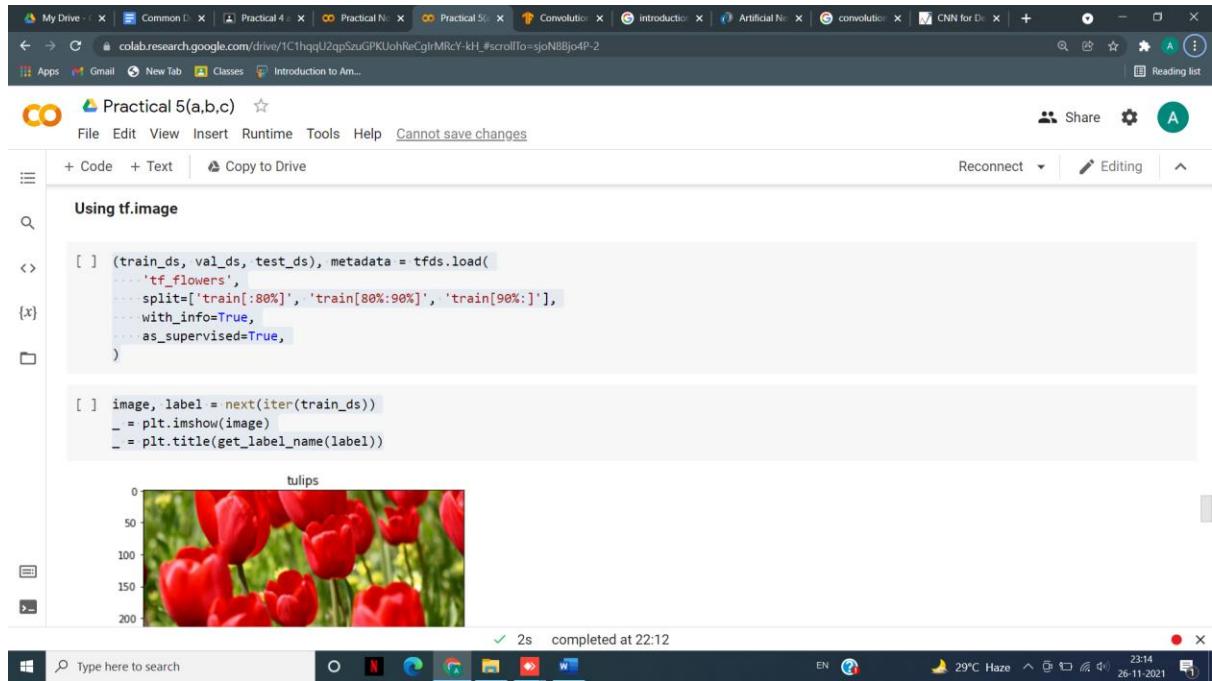
File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

Reconnect Editing

Using `tf.image`

```
<> [ ] (train_ds, val_ds, test_ds), metadata = tfds.load(  
{x}     ...'tf_flowers',  
     ...split=['train[:80%]', 'train[80%:90%]', 'train[90%:]'],  
     ...with_info=True,  
     ...as_supervised=True,  
)  
  
[ ] image, label = next(iter(train_ds))  
_ = plt.imshow(image)  
_ = plt.title(get_label_name(label))  
  
tulips  
0  
50  
100  
150  
200
```



My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural | convolution | CNN for D... | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

Reconnect Editing

```
seed = (i, 0) # tuple of size (2,)  
stateless_random_brightness = tf.image.stateless_random_brightness(  
    image, max_delta=0.95, seed=seed)  
visualize(image, stateless_random_brightness)
```

Original image Augmented image

Original image Augmented image

Original image Augmented image

2s completed at 22:12

Type here to search

EN 29°C Haze 23:15 26-11-2021

My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural | convolution | CNN for D... | +

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Apps Gmail New Tab Classes Introduction to Am...

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text

Reconnect Editing

Randomly crop an image

Randomly crop image using `tf.image.stateless_random_crop` by providing target size and seed. The portion that gets cropped out of `image` is at a randomly chosen offset and is associated with the given seed.

```
for i in range(3):  
    seed = (i, 0) # tuple of size (2,)  
    stateless_random_crop = tf.image.stateless_random_crop(  
        image, size=[210, 300, 3], seed=seed)  
    visualize(image, stateless_random_crop)
```

Original image Augmented image

Original image Augmented image

Original image Augmented image

Type here to search

EN 29°C Haze 23:15 26-11-2021

My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Digit Recognition | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] as_supervised=True,
[ ] 
```

```
[ ] def resize_and_rescale(image, label):
    image = tf.cast(image, tf.float32)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    image = (image / 255.0)
    return image, label
```

```
[ ] def augment(image_label, seed):
    image, label = image_label
    image, label = resize_and_rescale(image, label)
    image = tf.image.resize_with_crop_or_pad(image, IMG_SIZE + 6, IMG_SIZE + 6)
    # Make a new seed.
    new_seed = tf.random.experimental.stateless_split(seed, num=1)[0, :]
    # Random crop back to the original size.
    image = tf.image.stateless_random_crop(
        image, size=[IMG_SIZE, IMG_SIZE, 3], seed=new_seed)
    # Random brightness.
    image = tf.image.stateless_random_brightness(
        image, max_delta=0.5, seed=new_seed)
    image = tf.clip_by_value(image, 0, 1)
    return image, label
```

Type here to search

EN 29°C Haze 23:15 26-11-2021

My Drive | Practical 4 | Practical Notebooks | Practical 5 | Convolutional | introduction | Artificial Neural Networks | convolutional | CNN for Digit Recognition | + | Reading list

colab.research.google.com/drive/1C1hqqU2qpSzGPKUohReCgirMRcY-kH_#scrollTo=sjoN8Bjo4P-2

Practical 5(a,b,c)

File Edit View Insert Runtime Tools Help Cannot save changes

+ Code + Text Copy to Drive

```
[ ] 
```

```
[ ] 
```

```
[ ] def resize_and_rescale(image, label):
    image = tf.cast(image, tf.float32)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE])
    image = (image / 255.0)
    return image, label
```

Type here to search

EN 29°C Haze 23:16 26-11-2021

Practical No: 06

Aim: Building RNN.

What is RNN ?

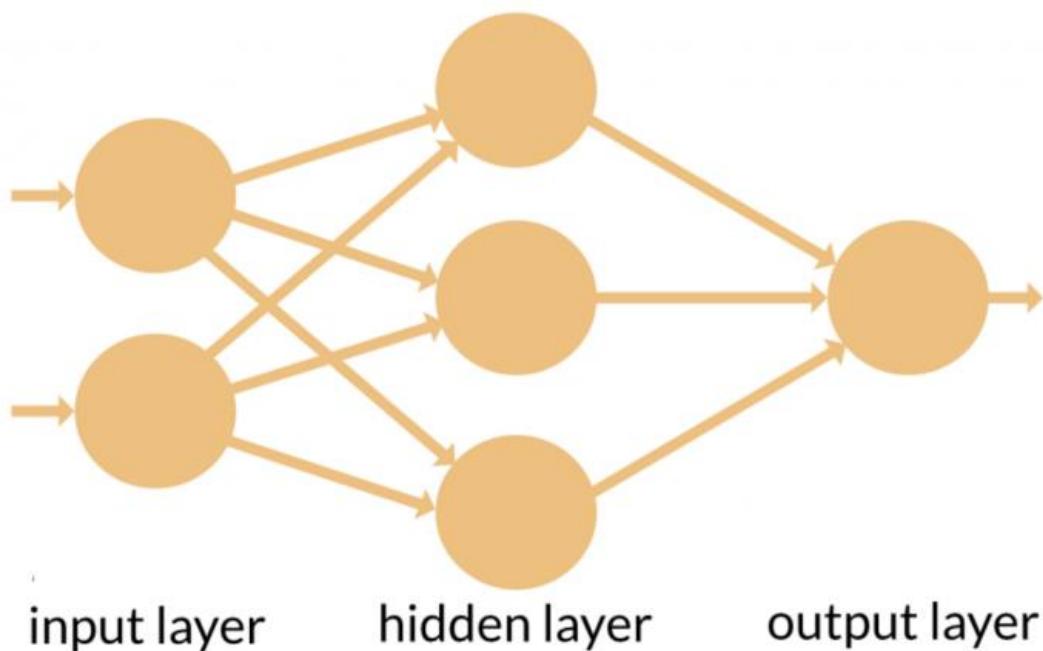
Ans. Recurrent neural networks (RNN) are the state of the art algorithm for sequential [data](#) and are used by Apple's Siri and Google's voice search. It is the first algorithm that remembers its input, due to an internal memory, which makes it perfectly suited for machine learning problems that involve sequential data. It is one of the algorithms behind the scenes of the amazing achievements seen in [deep learning](#) over the past few years.

How Recurrent Neural Networks Work

To understand RNNs properly, you'll need a working knowledge of "normal" feed-forward neural networks and sequential data.

Sequential data is basically just ordered data in which related things follow each other. Examples are financial data or the DNA sequence. The most popular type of sequential data is perhaps time series data, which is just a series of data points that are listed in time order.

RNN VS. FEED-FORWARD NEURAL NETWORKS



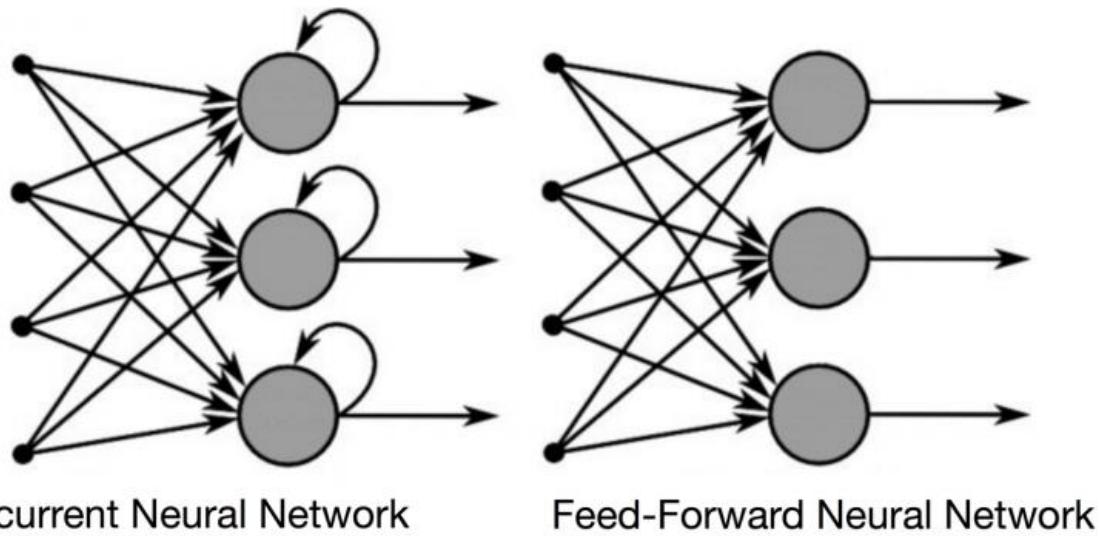
RNN's and feed-forward neural networks get their names from the way they channel information.

In a feed-forward neural network, the information only moves in one direction — from the input layer, through the hidden layers, to the output layer. The information moves straight through the network and never touches a node twice.

Feed-forward neural networks have no memory of the input they receive and are bad at predicting what's coming next. Because a feed-forward network only considers the current input, it has no notion of order in time. It simply can't remember anything about what happened in the past except its training.

In a RNN the information cycles through a loop. When it makes a decision, it considers the current input and also what it has learned from the inputs it received previously.

The two images below illustrate the difference in information flow between a RNN and a feed-forward neural network.



A usual RNN has a short-term memory. In combination with a LSTM they also have a long-term memory (more on that later).

Another good way to illustrate the concept of a recurrent neural network's memory is to explain it with an example:

Imagine you have a normal feed-forward neural network and give it the word "neuron" as an input and it processes the word character by character. By the time it reaches the character "r," it has already forgotten about "n," "e" and "u," which makes it almost impossible for this type of neural network to predict which character would come next.

A recurrent neural network, however, is able to remember those characters because of its internal memory. It produces output, copies that output and loops it back into the network.

Simply put: recurrent neural networks add the immediate past to the present.

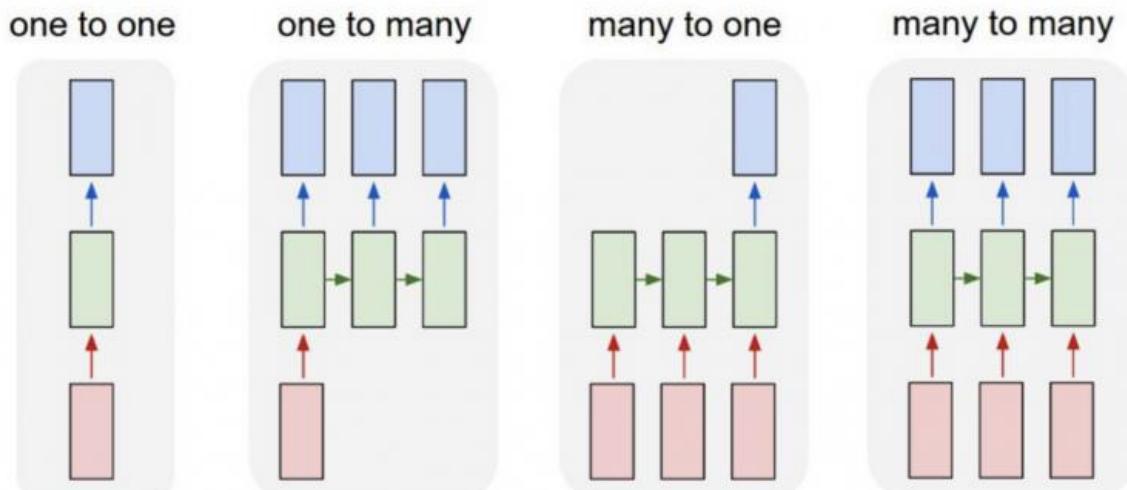
Therefore, a RNN has two inputs: the present and the recent past. This is important because the sequence of data contains crucial information about what is coming next, which is why a RNN can do things other algorithms can't.

A feed-forward neural network assigns, like all other deep learning algorithms, a weight matrix to its inputs and then produces the output. Note that RNNs apply weights to the current and also to the previous input. Furthermore, a recurrent neural network will also tweak the weights for both through gradient descent and backpropagation through time (BPTT).

TYPES OF RNNS

- One to One
- One to Many
- Many to One
- Many to Many

Also note that while feed-forward neural networks map one input to one output, RNNs can map one to many, many to many (translation) and many to one (classifying a voice).



Implementation of RNN:

https://colab.research.google.com/drive/1w9mit2ABNzawqSw1zLdySgf_tLIDF4qv

```
#setup
import numpy
import tensorflow
from tensorflow import keras
from tensorflow.keras import layers

#Here is a simple example of a Sequential model that processes sequences of integers, embeds each integer into a 64-dimensional vector, then processes the sequence of vectors using a LSTM layer.
model=keras.Sequential()
# Add an Embedding layer expecting input vocab of size 1000, and
# output embedding dimension of size 64.
model.add(layers.Embedding(input_dim=1000, output_dim=64))
#add a lstm layer with 128 internal units
model.add(layers.LSTM(128))
#add a dense layer with 10 units
model.add(layers.Dense(10))
model.summary()

model = keras.Sequential()
model.add(layers.Embedding(input_dim=1000, output_dim=64))

# The output of GRU will be a 3D tensor of shape (batch_size, timesteps, 256)
model.add(layers.GRU(256, return_sequences=True))

# The output of SimpleRNN will be a 2D tensor of shape (batch_size, 128)
model.add(layers.SimpleRNN(128))

model.add(layers.Dense(10))

model.summary()

encoder_vocab = 1000
decoder_vocab = 2000

encoder_input = layers.Input(shape=(None,))
encoder_embedded = layers.Embedding(input_dim=encoder_vocab, output_dim=64)(
```

```
    encoder_input
)

# Return states in addition to output
output, state_h, state_c = layers.LSTM(64, return_state=True, name="encoder")(
    encoder_embedded
)
encoder_state = [state_h, state_c]

decoder_input = layers.Input(shape=(None,))
decoder_embedded = layers.Embedding(input_dim=decoder_vocab, output_dim=64)(
    decoder_input
)

# Pass the 2 states to a new LSTM layer, as initial state
decoder_output = layers.LSTM(64, name="decoder")(
    decoder_embedded, initial_state=encoder_state
)
output = layers.Dense(10)(decoder_output)

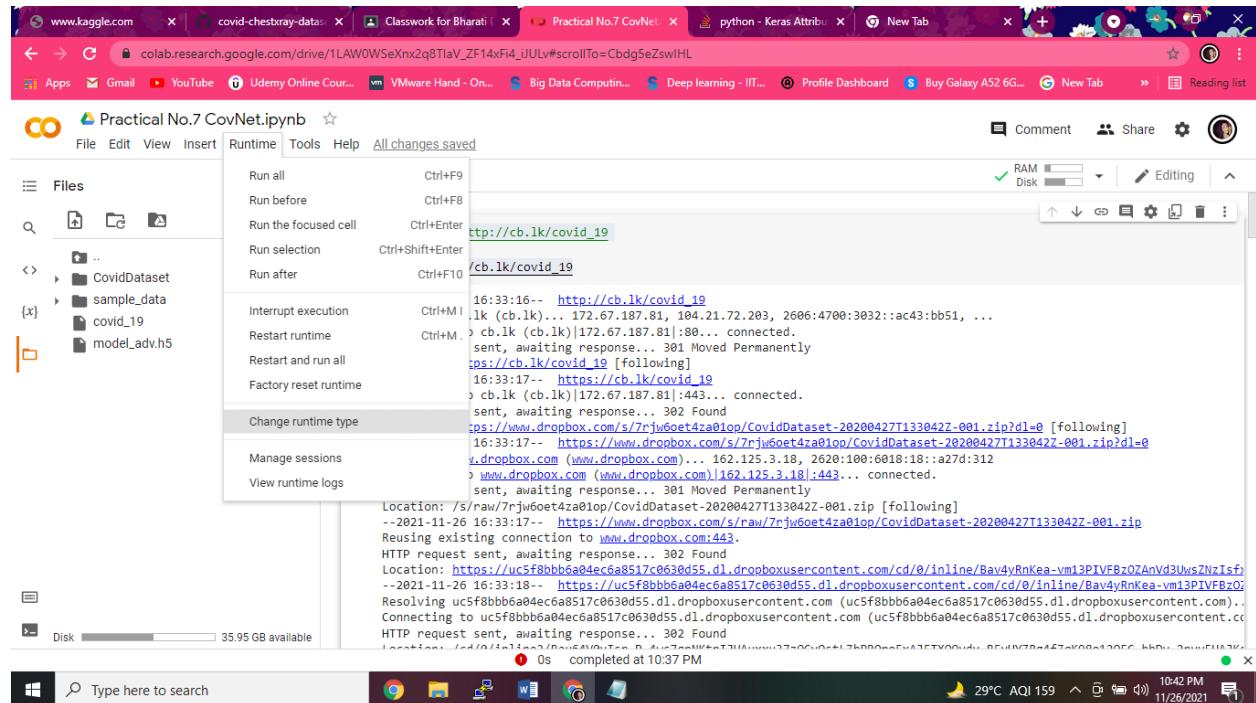
model = keras.Model([encoder_input, decoder_input], output)
model.summary()
```

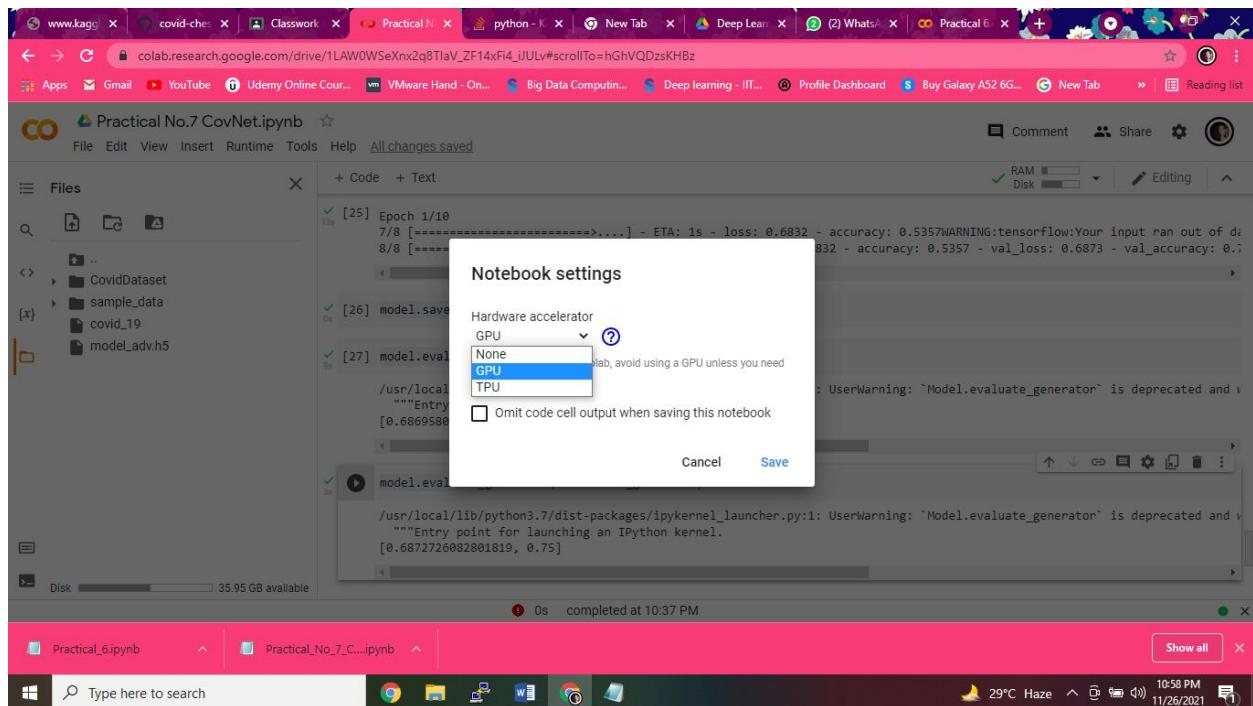
Practical No: 7

Aim: Build Deep Learning model using CovNet.

Implementation :

https://colab.research.google.com/drive/1LAW0WSeXnx2q8TlAV_ZF14xFi4_ijULv?usp=sharing





Code :

```
#Dataset : http://cb.lk/covid_19
```

```
!wget http://cb.lk/covid_19
```

```
!unzip covid_19
```

```
TRAIN_PATH = "CovidDataset/Train"  
VAL_PATH = "Covid_dataset/Test"
```

```
import numpy as np  
import matplotlib.pyplot as plt  
import keras  
from keras.layers import *  
from keras.models import *  
from keras.preprocessing import image
```

```
# CNN Based Model in Keras
```

```

model = Sequential()
model.add(Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=(224,224,3)))
model.add(Conv2D(64,(3,3),activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(64,(3,3),activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(128,(3,3),activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(64,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1,activation='sigmoid'))

model.compile(loss=keras.losses.binary_crossentropy,optimizer='adam',metrics=['accuracy'])

```

```

# Train from Scratch
train_datagen = image.ImageDataGenerator(
    rescale = 1./255,
    shear_range = 0.2,
    zoom_range = 0.2,
    horizontal_flip = True,
)
test_dataset = image.ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    'CovidDataset/Train',
    target_size = (224,224),
    batch_size = 32,
    class_mode = 'binary'
)

train_generator.class_indices

```

```
validation_generator = test_dataset.flow_from_directory(  
    'CovidDataset/Val',  
    target_size = (224,224),  
    batch_size = 32,  
    class_mode = 'binary'  
)
```

```
hist = model.fit_generator(  
    train_generator,  
    steps_per_epoch=8,  
    epochs = 10,  
    validation_data = validation_generator,  
    validation_steps=2  
)
```

```
model.save("model_adv.h5")
```

```
model.evaluate_generator(train_generator)
```

```
model.evaluate_generator(validation_generator)
```

Practical No: 8

Aim: Implementing a single Auto encoder based on fully connected layer

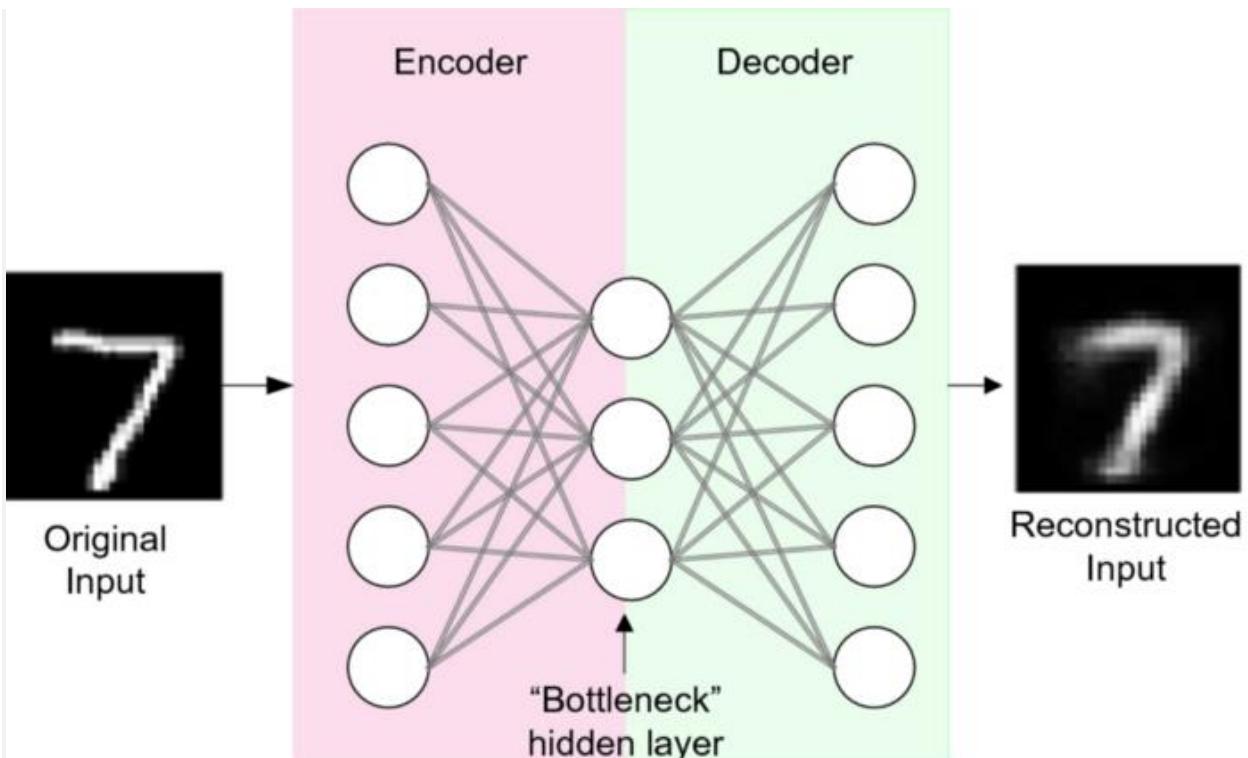
What is an Autoencoder ?

Neural networks exist in all shapes and sizes, and are often characterized by their **input** and **output data type**. For instance, image classifiers are built with *Convolutional Neural Networks*. They take **images** as inputs, and output a **probability distribution of the classes**.

Autoencoders (AE) are a family of neural networks for which **the input is the same as the output***. They work by compressing the input into a *latent-space representation*, and then reconstructing the output from this representation.

Auto-encoder is a **complex mathematical model** which trains on **unlabeled** as well as unclassified **data** and is used to map the **input data to** another compressed **feature representation** and from that feature representation **reconstructing back the input data**.

Auto-encoder is a **complex mathematical model** which trains on **unlabeled** as well as unclassified **data** and is used to map the **input data to** another compressed **feature representation** and from that feature representation **reconstructing back the input data**.



What is fully connected layer?

Fully Connected layers in a neural networks are those layers where all the inputs from one layer are connected to every **activation unit** of the next layer. In most popular machine learning models, the last few layers are full connected layers which **compiles the data extracted** by previous layers to form the final output. It is the second most time consuming layer second to Convolution Layer.

Code:

<https://colab.research.google.com/drive/1fNU9YBf5aYYEdhZJATBF5c22U-iZE2O?usp=sharing>

Single fully-connected neural layer as encoder and as decoder:

```
import keras
from keras import layers

# This is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats
```

```

# This is our input image
input_img = keras.Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = layers.Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = layers.Dense(784, activation='sigmoid')(encoded)

# This model maps an input to its reconstruction
autoencoder = keras.Model(input_img, decoded)

```

A separate encoder model:

```

# This model maps an input to its encoded representation
encoder = keras.Model(input_img, encoded)

```

decoder model:

```

# This is our encoded (32-dimensional) input
encoded_input = keras.Input(shape=(encoding_dim,))
# Retrieve the last layer of the autoencoder model
decoder_layer = autoencoder.layers[-1]
# Create the decoder model
decoder = keras.Model(encoded_input, decoder_layer(encoded_input))

```

Train our autoencoder to reconstruct MNIST digits.

we'll configure our model to use a per-pixel binary crossentropy loss, and the Adam optimizer:

```
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

Let's prepare our input data. We're using MNIST digits, and we're discarding the labels (since we're only interested in encoding/decoding the input images).

```

from keras.datasets import mnist
import numpy as np
(x_train, _), (x_test, _) = mnist.load_data()

```

We will normalize all values between 0 and 1 and we will flatten the 28x28 images into vectors of size 784.

```

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
print(x_train.shape)

```

```
print(x_test.shape)
```

Now let's train our autoencoder for 50 epochs:

```
autoencoder.fit(x_train, x_train,
                 epochs=50,
                 batch_size=256,
                 shuffle=True,
                 validation_data=(x_test, x_test))
```

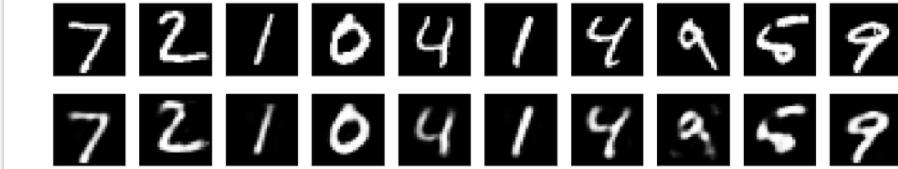
After 50 epochs, the autoencoder seems to reach a stable train/validation loss value of about 0.09. We can try to visualize the reconstructed inputs and the encoded representations. We will use Matplotlib.

```
# Encode and decode some digits
# Note that we take them from the *test* set
encoded_imgs = encoder.predict(x_test)
decoded_imgs = decoder.predict(encoded_imgs)

# Use Matplotlib (don't ask)
import matplotlib.pyplot as plt

n = 10 # How many digits we will display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



Practical No.8 Autoencoder.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

Connecting to a runtime to enable file browsing.

```
# Display reconstruction
... ax = plt.subplot(2, n, i + 1 + n)
... plt.imshow(decoded_imgs[i].reshape(28, 28))
... plt.gray()
... ax.get_xaxis().set_visible(False)
... ax.get_yaxis().set_visible(False)
plt.show()
```

Reconnect | Editing

Here's what we get. The top row is the original digits, and the bottom row is the reconstructed digits. We are losing quite a bit of detail with this basic approach.

Type here to search

Google Colab Link :-

<https://drive.google.com/drive/folders/1raAACU-jPioJy8CNO9OwvaYWILijWXQ?usp=sharing>

